

# AgentAlign Lab: Master Execution Handbook

V3 final execution edition: daily execution system, debugging companion, implementation navigator, experimental plan, and internship positioning playbook. The original 53-section handbook is preserved in full; Sections 54-70 add the operational layer for building the project over 6-8 weeks, and the early unnumbered sections clarify research questions, success criteria, non-goals, limitations, ablations, and evaluation integrity.

Version: 2026-05-23

Purpose: a brutally practical, internship-optimized AI systems roadmap for building a verifier-guided DPO project with RLAIIF-style feedback around terminal agents.

Audience: solo 3rd-year AIML student with strong ML/DL and Hugging Face skills, basic terminal comfort, beginner infra/Docker/observability, free GPU access, and a 6 to 8 week timeline.

Project target: AI/ML internship proof-of-work for YC AI startups, Bangalore AI startups, AI infra companies, evals/observability companies, post-training/alignment startups, and agent/tool-use companies.

Important principle: this is not a giant academic benchmark project. It is a compact, shipped AI systems project with measurable results, a clear demo, and a strong technical story.

---

## 1. Executive Summary

AgentAlign Lab is a compact AI systems project that builds a small terminal-agent evaluation and post-training pipeline.

The project creates deterministic coding/data tasks, runs a small open-source LLM as a terminal agent, records every action as a structured trajectory, scores the trajectory with verifiers, converts better vs worse trajectories into preference pairs, and fine-tunes the model with DPO/LoRA. The final output is a dashboard and report showing whether verifier-guided preference learning improves tool-use reliability.

The key portfolio signal is:

I can build the feedback loop that modern agent companies need: run agents, trace actions, verify outcomes, create preference data, post-train models, and evaluate failures.

The final MVP should produce:

- 60 to 100 terminal tasks.
- 300 to 700 agent trajectories.
- 300 to 700 preference pairs, with 800 to 1,500 as a strong stretch target.
- One DPO/LoRA fine-tuned small model.
- Baseline vs tuned comparison.
- A Gradio dashboard.

- A strong README.
- A 6 to 8 page paper-style technical report.
- A 60 to 90 second demo video.

The project should not try to become SWE-bench, Terminal-Bench, LangSmith, or Rubric. It should instead become a small but complete slice of the same problem space.

```
flowchart LR
  A["Small terminal tasks"] --> B["Agent runs"]
  B --> C["Trajectory logs"]
  C --> D["Deterministic verifiers"]
  D --> E["Preference pairs"]
  E --> F["DPO fine-tuning"]
  F --> G["Evaluation"]
  G --> H["Dashboard + report + outreach"]
```

---

## Research Questions

This project is not only an implementation exercise. It is a compact experiment about whether deterministic execution feedback can improve terminal-agent reliability.

Primary research question:

Can deterministic verifier feedback improve terminal-agent reliability through DPO preference optimization?

Secondary questions:

- Does verifier-guided DPO reduce unsafe tool usage?
- Does verifier-guided DPO reduce invalid tool actions?
- Does DPO improve efficiency and step quality, not only final pass rate?
- Which task families benefit most from verifier-generated preferences?
- Do high-margin preference pairs outperform noisy or low-margin pairs?

Why this matters:

Modern AI startups increasingly build feedback loops around agents: run the agent, trace behavior, score outcomes, improve prompts/data/models, and evaluate regressions. AgentAlign Lab studies that loop at a solo-student scale. The goal is not to prove a universal alignment result. The goal is to produce credible local evidence about a narrow, measurable system.

---

## MVP Success Criteria

The MVP is successful if all of the following are true:

- 60+ verified tasks exist.
- 300+ trajectories are collected.

- DPO training completes successfully.
- Held-out evaluation runs reproducibly.
- Dashboard visualizes traces and metrics.
- At least one held-out metric improves after tuning.

The metric that improves can be:

- held-out pass rate
- invalid action rate
- unsafe action rate
- timeout rate
- average steps on solved tasks
- verifier score
- failure rate for a specific high-priority failure tag

Important:

The MVP does not require every metric to improve. It requires a complete loop, a reproducible evaluation, and at least one credible held-out improvement.

This prevents two common traps:

- endless tuning because pass rate is not perfect
- scope creep because the project feels "not researchy enough"

## Non-Goals

AgentAlign Lab is intentionally constrained. These non-goals are part of the engineering discipline, not a weakness.

| Non-goal                         | Why it is excluded                                                   |
|----------------------------------|----------------------------------------------------------------------|
| Frontier RLHF                    | Too expensive, unstable, and unnecessary for internship signal       |
| PPO/RLHF implementation          | DPO is simpler, cheaper, and enough for verifier-derived preferences |
| Large-scale distributed training | Free-GPU feasibility is a core constraint                            |
| Autonomous SWE agent benchmark   | Real SWE tasks add setup complexity and benchmark contamination risk |
| Production deployment            | The goal is a reproducible local AI systems artifact, not a SaaS     |
| Multi-agent orchestration        | Adds complexity without improving the core feedback loop             |
| Browser automation               | Interesting, but expands task surface too much                       |

| Non-goal                             | Why it is excluded                                                          |
|--------------------------------------|-----------------------------------------------------------------------------|
|                                      | for MVP                                                                     |
| Large benchmark replication          | The project should be a clean small eval, not a half-built public benchmark |
| Human preference collection platform | Deterministic verifiers are the supervision source for MVP                  |
| Reward model training                | Useful future work, but DPO from verifier-ranked traces is enough           |
| MCP tool server                      | Important modern standard, but V3 only after the local tool layer works     |
| Full observability platform clone    | Local traces and dashboard first; Phoenix/Langfuse are V2 integrations      |

The core discipline:

If a feature does not improve task quality, trace quality, preference quality, evaluation credibility, or demo clarity, it is not MVP.

## Known Limitations

These limitations should be stated openly in the README and technical report. They make the project more credible because they show evaluation awareness.

| Limitation                                                | Why it matters                                                        | Mitigation                                                           |
|-----------------------------------------------------------|-----------------------------------------------------------------------|----------------------------------------------------------------------|
| Deterministic verifiers can be gamed                      | Agents may learn shortcuts that satisfy tests without robust behavior | protect test files, add hidden checks, inspect diffs, audit failures |
| Small-model findings may not generalize                   | A 0.5B-1.5B model behaves differently from frontier models            | frame results as local systems evidence, not universal claims        |
| Local tasks are narrower than real SWE work               | Real repositories have dependencies, ambiguity, and long context      | keep claims focused on small terminal tasks                          |
| Verifier-guided optimization can overfit heuristics       | DPO may reward verifier artifacts                                     | use held-out tasks, high-margin pairs, failure analysis              |
| Terminal tasks do not fully measure long-horizon planning | Max 8-step tasks are intentionally compact                            | call this an MVP constraint and future direction                     |
| Preference quality depends heavily on verifier quality    | Bad verifiers create bad chosen/rejected data                         | reference solution checks, broken-template checks, manual audits     |
| DPO may improve format reliability more than task success | Structured action gains may be larger than pass-rate gains            | report multiple metrics, not only pass rate                          |
| Generated tasks may contain                               | Synthetic scale can introduce                                         | validate every task and keep                                         |

| Limitation  | Why it matters | Mitigation        |
|-------------|----------------|-------------------|
| hidden bugs | noise          | manual seed tasks |

How to phrase this in interviews:

The project is intentionally narrow. I optimized for a reproducible feedback loop, not a broad benchmark claim. The limitations are exactly why I track verifier quality, held-out splits, and failure modes.

## Future Research Directions

These ideas are intentionally excluded from MVP scope. They are useful for interviews and V2/V3 planning, but they should not delay the 6-8 week build.

| Direction                           | Why it is interesting                                             | MVP status  |
|-------------------------------------|-------------------------------------------------------------------|-------------|
| Process supervision                 | Scores intermediate reasoning/tool-use steps                      | excluded    |
| Step-level DPO                      | Trains on local action quality, not only final trajectory outcome | excluded    |
| AI critique/revision loops          | Uses model critics to improve trajectories or labels              | excluded    |
| Reward models                       | Learns a reward function beyond deterministic verifiers           | excluded    |
| Online RL methods such as GRPO/RLOO | Optimizes through interaction with reward signals                 | excluded    |
| MCP-compatible tool interfaces      | Connects the system to modern agent-tool standards                | excluded    |
| Browser agents                      | Expands from terminal tasks to web workflows                      | excluded    |
| Long-horizon tasks                  | Tests planning over many steps                                    | excluded    |
| Multi-agent traces                  | Studies collaboration/debate/tool delegation                      | excluded    |
| Dockerized reproducible evaluation  | Improves isolation and portability                                | V2, not MVP |

The right future-work sentence:

After the MVP is demoable, the highest-value next step is Dockerized reproducible evaluation plus richer trace export, not jumping directly into online RL.

## Planned Ablation Experiments

These ablations make the project experimentally grounded without turning it into a research lab. Run only the first one or two if time is tight.

Primary metric:

held-out pass rate

Secondary metrics:

- invalid action rate
- unsafe action rate
- timeout rate
- average steps on solved tasks
- verifier score
- failure-tag distribution
- pair rejection/audit rate

Evaluation methodology:

- split tasks before trajectory generation
- keep held-out task IDs fixed
- run baseline and tuned models with the same max steps, tools, prompt template, and verifier
- compare each ablation on the same held-out evaluation split
- report denominator, config, dataset version, and random seed where applicable

| Ablation                                       | Comparison                                        | Hypothesis                                                     | Primary metric        | Secondary metrics                 |
|------------------------------------------------|---------------------------------------------------|----------------------------------------------------------------|-----------------------|-----------------------------------|
| High-margin vs all pairs                       | train on score-margin-filtered pairs vs all pairs | cleaner preferences improve reliability more than noisy volume | held-out pass rate    | invalid action rate, failure tags |
| Safety penalties vs none                       | include unsafe-action rejections vs exclude them  | explicit safety contrast reduces forbidden behavior            | unsafe action rate    | pass rate, timeout rate           |
| Bug-fix-only vs mixed-task training            | train on Python bug-fix pairs vs all families     | narrow training may help one family but generalize less        | family-wise pass rate | invalid action rate               |
| Low-temperature vs high-temperature trajectory | sample conservative vs diverse attempts           | diversity creates better rejected examples but more noise      | held-out pass rate    | pair quality, timeout rate        |

| Ablation                                          | Comparison                                          | Hypothesis                                       | Primary metric      | Secondary metrics    |
|---------------------------------------------------|-----------------------------------------------------|--------------------------------------------------|---------------------|----------------------|
| generation                                        |                                                     |                                                  |                     |                      |
| Trajectory-level vs final-answer-only preferences | train on full action traces vs final responses only | full trajectories teach tool-use behavior better | invalid action rate | avg steps, pass rate |

Minimum viable ablation:

High-margin pairs vs all pairs.

If only one ablation is completed, make it that one.

## Evaluation Integrity

Evaluation integrity is what separates a credible AI systems project from a demo that accidentally fools itself.

Rules:

- Split tasks before trajectory generation.
- Never generate training preferences from held-out tasks.
- Use deterministic verifiers for MVP scoring.
- Run broken-template and reference-solution checks for each task family.
- Add anti-cheat protections for tests, expected outputs, and forbidden file edits.
- Manually audit a sample of preference pairs before DPO.
- Keep train, validation, and held-out test task IDs separate.
- Save config, dataset version, model name, adapter path, and command for every experiment.
- Report limitations and failure modes with the metric table.

Anti-contamination flow:

```
graph TD
    A["Task suite"] --> B["Split by task_id"]
    B --> C["Train tasks"]
    B --> D["Validation tasks"]
    B --> E["Held-out test tasks"]
    C --> F["Generate train trajectories"]
    F --> G["Build DPO preference pairs"]
    G --> H["Train adapter"]
    D --> I["Tune configs and debug"]
    E --> J["Final held-out evaluation only"]
    H --> J
```

What not to do:

- do not split individual trajectories randomly after all runs are collected

- do not tune prompts repeatedly on the final held-out set
- do not report the best cherry-picked run
- do not train on tasks that appear in the final dashboard metrics

Founder-readable version:

I split by task before generating trajectories, so the DPO dataset cannot contain behavior from final held-out tasks. The final metrics are small-scale but not contaminated by train/test overlap.

---

## System Design Tradeoff Table

This table is the systems-engineering backbone of the project. It explains why the MVP is compact by design.

| Decision            | Chosen approach                 | Tradeoff                                                             |
|---------------------|---------------------------------|----------------------------------------------------------------------|
| Optimization method | DPO instead of PPO              | simpler and more stable, less RL realism                             |
| Feedback source     | deterministic verifiers         | reliable scoring, narrower behavior coverage                         |
| Execution isolation | local tempdir sandbox           | faster iteration, weaker isolation than Docker                       |
| Model size          | 0.5B-1.5B code model            | affordable/free, lower capability ceiling                            |
| Agent architecture  | custom single-agent loop        | better systems understanding, less framework abstraction             |
| Tool format         | strict JSON actions             | easier validation, less natural flexibility                          |
| Observability       | local JSONL traces + Gradio     | reproducible and simple, less production-grade than Phoenix/Langfuse |
| Dataset scale       | 60+ tasks and 300+ trajectories | buildable solo, not benchmark-scale                                  |
| Task source         | small local tasks               | controlled verification, less real-world SWE breadth                 |
| Safety scope        | allowlisted local commands      | practical safety, not production-grade sandbox security              |

The key message:


Every simplification is chosen to preserve finishability, reproducibility, and measurement quality.

---




# Data Flywheel

AgentAlign is valuable because it builds a feedback loop, not because it runs one model once.




Artifact view:



Why this matters:

This mirrors the operating loop inside modern post-training, evals, and agent-infrastructure startups:



The MVP implements this loop locally and visibly.

## Engineering Learning Outcomes

By finishing the MVP, you will be able to credibly discuss and demonstrate:

| Capability                    | Project artifact                                  |
|-------------------------------|---------------------------------------------------|
| Structured tool calling       | JSON action schema + parser                       |
| Verifier-based evaluation     | pytest/schema/output/safety verifiers             |
| Trajectory logging            | step-by-step trace JSONL                          |
| Preference dataset generation | chosen/rejected trajectory pairs                  |
| DPO pipelines                 | TRL DPOTrainer script and experiment report       |
| LoRA/QLoRA training           | small-model adapter training on free GPU          |
| Observability concepts        | local trace dashboard and failure tables          |
| Sandboxed execution           | tempdir workspace + allowlisted commands          |
| Failure analysis              | failure taxonomy and before/after breakdown       |
| Agent reliability engineering | invalid action, unsafe action, timeout, and pass- |

| Capability | Project artifact |
|------------|------------------|
|            | rate metrics     |

Recruiter translation:

This project maps to AI infra, evals, observability, post-training, and agent-tooling internships because it shows the whole feedback loop, not only a model call.

## Glossary

| Term            | Meaning                                                                                              |
|-----------------|------------------------------------------------------------------------------------------------------|
| trajectory      | sequence of agent prompts, actions, observations, errors, and final verifier result                  |
| verifier        | deterministic correctness/safety checker that scores a task attempt                                  |
| preference pair | chosen vs rejected trajectory for the same task prompt                                               |
| DPO             | Direct Preference Optimization, a preference-learning method that trains on chosen/rejected examples |
| LoRA            | low-rank fine-tuning technique that trains small adapter weights                                     |
| QLoRA           | quantized LoRA fine-tuning for lower GPU memory usage                                                |
| unsafe action   | forbidden tool behavior such as deleting tests, accessing network, or writing outside workspace      |
| invalid action  | model output that cannot be parsed or validated as an allowed tool action                            |
| held-out task   | task never used for training preference data                                                         |
| failure tag     | structured label explaining why a trajectory failed                                                  |

## Checkpoint Milestones

Use these as psychological and technical checkpoints.

| Milestone | What it means         | Evidence                                   |
|-----------|-----------------------|--------------------------------------------|
| 1         | verifier works        | broken tasks fail and reference fixes pass |
| 2         | agent completes tasks | baseline agent attempts tasks              |

|  | Milestone | What it means                     | Evidence                                          |
|--|-----------|-----------------------------------|---------------------------------------------------|
|  |           |                                   | and saves results                                 |
|  | 3         | trajectories collected            | 300+ structured trajectories exist                |
|  | 4         | DPO training succeeds             | adapter trains and loads                          |
|  | 5         | evaluation shows measurable delta | at least one held-out metric improves             |
|  | 6         | dashboard + demo complete         | traces, metrics, and failure analysis are visible |

Do not skip milestones. If Milestone 1 is weak, every later result becomes suspect.

---

## Kill Criteria

This project should not become endless tuning hell.

If DPO does not produce measurable improvement after:

- verifier auditing
- high-margin pair filtering
- train/eval leakage checks
- at least two small DPO runs with sane hyperparameters
- inspection of before/after model outputs

then stop chasing training gains and pivot the project emphasis toward:

- evaluation infrastructure
- observability
- trajectory analysis
- verifier engineering
- preference dataset quality
- failure taxonomy

This is still a strong internship project.

The mature framing:

The training result was mixed, so I turned the project into a rigorous agent evaluation and trace-analysis system. That is a realistic AI engineering outcome, not a failure.

The worst outcome is not "DPO did not improve pass rate." The worst outcome is spending four extra weeks tuning without learning anything or shipping a clear demo.

---

## 2. Final Revised Project Vision

Final project title:

**AgentAlign Lab: Verifier-Guided Preference Learning for Reliable Terminal Agents**

Final project subtitle:

**A compact post-training and observability pipeline that turns terminal-agent traces into DPO data.**

The revised project is not:

- a huge RLHF system
- a full coding-agent benchmark
- a Docker-first sandbox platform
- a browser-agent project
- a multi-agent framework
- a full observability SaaS

The revised project is:

- a small local agent harness
- a task and verifier suite
- a trace logging system
- a preference data generator
- a DPO training pipeline
- an evaluation dashboard
- a failure analysis report

The core contribution:

Instead of asking "Can an LLM solve this task?", AgentAlign Lab asks "Can we use deterministic feedback from task execution to generate preference data that improves future agent behavior?"

This is high leverage because startups care deeply about:

- agent reliability
- evals
- trace analysis
- post-training
- tool-use safety
- measurable product improvement
- small models getting better on narrow workflows

---

### 3. Why This Project Is Strong for Internships

This project is strong because it demonstrates multiple startup-relevant skills in one coherent artifact.

| Skill signal     | What you build                                         | Why startups care                             |
|------------------|--------------------------------------------------------|-----------------------------------------------|
| ML engineering   | DPO/LoRA fine-tuning on small models                   | Shows practical post-training ability         |
| AI systems       | Agent loop, tools, logs, verifiers                     | Shows you can build more than notebooks       |
| Evals            | Held-out tasks, success metrics, failure taxonomy      | Directly relevant to eval/observability teams |
| Agent infra      | Tool schema, command execution, retries, safety checks | Relevant to agent/tool-use startups           |
| Data engineering | Trajectories to preference pairs                       | Relevant to data and post-training companies  |
| Product demo     | Dashboard, README, video                               | Recruiter/founder readable                    |
| Research taste   | Ablations and limitations                              | Shows maturity without overdoing theory       |

Founder mental reaction you want:

This student understands the modern AI stack. They can build evals, collect traces, make preference data, fine-tune small models, and explain failures clearly.

That is much stronger than:

- "I built a chatbot."
- "I trained CNNs on CIFAR."
- "I made a RAG app."
- "I read RLHF papers."

---

### 4. What Recruiters/Founders Actually Care About

Early-stage AI founders usually do not inspect every line of code first. They scan for signal.

They care about:

1. Does the project solve a real problem they recognize?
2. Is there a working demo?
3. Are the results measurable?
4. Is the README clear in 60 seconds?

5. Did the student avoid obvious overengineering?
6. Can the student explain tradeoffs?
7. Does the student understand failure modes?
8. Can this person ship without being handheld?

They do not care much about:

- a massive theoretical literature review
- a huge but broken benchmark suite
- an overcomplicated microservice architecture
- fancy diagrams with no working code
- vague claims like "improves alignment"
- chasing public leaderboard numbers

Recruiter-readable proof:

Built a verifier-guided DPO pipeline for terminal agents:  
100 tasks, 600 trajectories, 1,200 preference pairs, Qwen2.5-Coder-1.5B LoRA  
fine-tune,  
baseline vs tuned eval, Gradio observability dashboard, and failure analysis.

Founder-readable proof:

I built a mini version of the loop agent-infra companies need:  
run agents -> trace behavior -> verify outcomes -> generate feedback -> post-  
train -> evaluate regressions.

---

## 5. MVP vs V2 Breakdown

The biggest danger is trying to build V2 before MVP. Do not do that.

### MVP

Build this in 6 weeks:

- Local Python task runner.
- 60 to 100 tasks.
- Mostly handwritten tasks, with optional LLM-assisted task drafts only after verifier validation works.
- Minimal ReAct-style agent loop.
- Structured JSON actions.
- Deterministic verifier.
- Trajectory logs.
- Preference pair builder.

- DPO/LoRA training.
- Baseline vs tuned evaluation.
- Gradio dashboard.
- README + report.

V2

Add after the MVP works:

- Docker sandbox.
- Step-level preference pairs.
- AI judge critique.
- Phoenix or Langfuse tracing.
- Hugging Face dataset/model upload.
- Optional LangGraph reimplementation.
- Optional MCP-style tool interface.
- Optional API server.

Do not invert the order

```
flowchart TD
  A["MVP: local harness"] --> B["MVP: deterministic verifier"]
  B --> C["MVP: preference data"]
  C --> D["MVP: DPO"]
  D --> E["MVP: dashboard"]
  E --> F["V2: Docker"]
  F --> G["V2: observability platform"]
  G --> H["V2: LangGraph/MCP/API"]
```

6. Exact Final MVP Scope

Final MVP name:

AgentAlign Lab MVP

MVP promise:

Given a suite of small terminal tasks, compare a baseline model and a DPO-tuned model on verified tool-use behavior.

MVP feature list

| Component  | MVP requirement                                |
|------------|------------------------------------------------|
| Task suite | 60 to 100 tasks across 5 compact task families |
| Agent loop | single-agent, JSON actions, max 8 steps        |

| Component    | MVP requirement                                                      |
|--------------|----------------------------------------------------------------------|
| Tools        | read file, write file, list files, run allowed command, final answer |
| Sandbox      | Python tempdir with allowlisted subprocess commands                  |
| Verifier     | pytest/schema/output/safety checks                                   |
| Trajectories | JSONL logs with every step                                           |
| Preferences  | chosen/rejected pairs from verifier ranking                          |
| Training     | QLoRA/DPO on 0.5B to 1.5B model                                      |
| Eval         | held-out tasks, baseline vs tuned                                    |
| Dashboard    | Gradio, local only                                                   |
| Report       | concise paper-style writeup                                          |

## MVP task families

| Task family         | Count | Verifier                  |
|---------------------|-------|---------------------------|
| Python bug fix      | 30    | pytest -q                 |
| Data transformation | 25    | exact output + schema     |
| JSON/config repair  | 20    | parser + expected fields  |
| Text/log extraction | 15    | regex/exact output        |
| Safety traps        | 10    | forbidden action detector |

Total: 100 tasks.

If overwhelmed, cut to 60 tasks:

- 25 Python bug fix
- 15 data transformation
- 10 JSON/config repair
- 5 text extraction
- 5 safety traps

## 7. What To REMOVE Completely

Remove these from the first 6 weeks.

| Remove        | Why                                                       |
|---------------|-----------------------------------------------------------|
| Full PPO RLHF | Too expensive, unstable, not needed for internship signal |



| Remove                              | Why                                     |
|-------------------------------------|-----------------------------------------|
| Large 7B+ training                  | Free GPU bottleneck, slow iteration     |
| Browser automation                  | Hard to verify, high infra cost         |
| Multi-agent orchestration           | Distracts from post-training loop       |
| Firecracker/microVMs                | Great infra, too much for this project  |
| Real cybersecurity tasks            | Safety/legal complexity                 |
| Huge benchmark replication          | Low novelty, high setup burden          |
| Kubernetes                          | Irrelevant for solo MVP                 |
| Complex backend/frontend            | Dashboard is enough                     |
| MCP implementation                  | Useful concept, not needed for MVP      |
| LangGraph as core runtime           | Study it, but build your own loop first |
| Phoenix/Langfuse as hard dependency | V2 only; local logs first               |

Hard rule:

If a feature does not improve demo clarity, measured results, or learning value within 6 weeks, cut it.

## 8. What To DEFER To V2

V2 features are valuable only after the MVP works.

| V2 feature           | Why useful later                                    | When to add                        |
|----------------------|-----------------------------------------------------|------------------------------------|
| Docker sandbox       | Stronger safety/reproducibility                     | After local runner works           |
| Step-level DPO       | More research-grade feedback                        | After trajectory-level DPO works   |
| AI judge             | Adds AI-assisted critique and RLAIIF-style analysis | After deterministic verifier works |
| Phoenix tracing      | Modern observability signal                         | After logs are stable              |
| Langfuse tracing     | Good for LLM app observability                      | After dashboard exists             |
| LangGraph version    | Shows industry framework awareness                  | After custom loop is complete      |
| Hugging Face dataset | Public artifact for portfolio                       | After dataset is clean             |
| Model card           | Professional ML release skill                       | After training results are stable  |
| API server           | Product polish                                      | After demo works locally           |

V2 success condition:

The MVP is already demoable, and V2 makes it more credible, not merely bigger.

---

## 9. Final Simplified Architecture

flowchart TD

```
A["Task JSONL"] --> B["Task Loader"]
B --> C["Temp Workspace Builder"]
C --> D["Agent Loop"]
E["Base Model"] --> D
D --> F["Tool Executor"]
F --> D
D --> G["Trajectory Logger"]
G --> H["Verifier"]
H --> I["Scored Trajectories"]
I --> J["Preference Builder"]
J --> K["DPO Dataset"]
K --> L["DPO Trainer"]
L --> M["LoRA Adapter"]
M --> N["Evaluation Runner"]
E --> N
N --> O["Metrics JSON/CSV"]
O --> P["Gradio Dashboard"]
O --> Q["Report"]
```

Key simplification:

- No external services for MVP.
- No Docker for MVP.
- No framework dependency for agent loop.
- All artifacts are local JSONL/CSV/Markdown.

This makes the project buildable and inspectable.

---

## 10. End-to-End System Workflow

### System lifecycle

sequenceDiagram

```
participant Task as Task Suite
participant Runner as Agent Runner
participant Model as LLM
participant Tools as Local Tools
participant Log as Trajectory Log
participant Verifier as Verifier
participant Pref as Preference Builder
participant Train as DPO Trainer
participant Eval as Evaluator
participant UI as Dashboard

Task->>Runner: Load task files and instruction
Runner->>Model: Prompt with tool schema and history
Model->>Runner: JSON action
Runner->>Tools: Execute safe action
Tools->>Runner: Observation
Runner->>Log: Append step
Runner->>Model: Continue until final or max steps
```

```
Runner->>Verifier: Run tests/checks
Verifier->>Log: Attach score
Log->>Pref: Rank trajectories
Pref->>Train: DPO chosen/rejected data
Train->>Eval: Tuned adapter
Eval->>UI: Baseline vs tuned metrics
```

## Artifact flow

```
flowchart LR
    A["data/tasks/*.jsonl"] --> B["runs/*.jsonl"]
    B --> C["data/trajectories/*.jsonl"]
    C --> D["data/preferences/*.jsonl"]
    D --> E["outputs/adapters/*"]
    E --> F["results/*.json"]
    F --> G["dashboard"]
    F --> H["report/paper.md"]
```

## 11. Recommended Modern Tech Stack (2025-2026)

### Final MVP stack

| Layer               | Tool                                   | Why                                     |
|---------------------|----------------------------------------|-----------------------------------------|
| Language            | Python 3.11+                           | Best ecosystem for ML and agent tooling |
| Package manager     | uv or pip + venv                       | uv is fast and modern; pip is okay      |
| Data validation     | Pydantic v2                            | Great for task/action schemas           |
| LLM inference       | Hugging Face Transformers              | Simple for local/Colab small models     |
| Fine-tuning         | TRL + PEFT + bitsandbytes              | Standard DPO/LoRA stack                 |
| Dataset handling    | Hugging Face Datasets + JSONL          | Simple and portable                     |
| Verification        | pytest + custom Python checks          | Deterministic and startup-readable      |
| Dashboard           | Gradio Blocks                          | Fast ML demo UI                         |
| Experiment tracking | CSV/JSON first, W&B or MLflow optional | Avoid setup drag                        |
| Observability       | local trace JSON first                 | Full platforms are V2                   |
| Sandbox             | Python tempdir + subprocess allowlist  | Buildable before Docker                 |

## V2 stack

| Layer               | Tool                                | Use                                     |
|---------------------|-------------------------------------|-----------------------------------------|
| Sandbox             | Docker                              | Stronger isolated task execution        |
| Observability       | Phoenix or Langfuse                 | Trace UI, spans, evals                  |
| Agent framework     | LangGraph or OpenAI Agents SDK      | Study and optionally reimplement        |
| Model serving       | vLLM, SGLang, Ollama, or llama.cpp  | Only if you deploy or compare inference |
| Experiment tracking | W&B or MLflow                       | If running many configs                 |
| Eval framework      | Inspect AI or Braintrust-style Eval | If formalizing eval runs                |

## What is outdated or lower signal

| Tool/pattern                                | Status                                 | Advice                                 |
|---------------------------------------------|----------------------------------------|----------------------------------------|
| Raw LangChain chains for everything         | Often too abstract/noisy               | Learn LangGraph patterns instead       |
| CrewAI for this MVP                         | Too productized for your core learning | Skip                                   |
| AutoGen for this MVP                        | Multi-agent complexity too early       | Skip                                   |
| PPO RLHF tutorials as implementation target | Too heavy                              | Understand conceptually, implement DPO |
| Notebook-only project                       | Weak startup signal                    | Use scripts, package structure, CLI    |
| Streamlit-only metrics app                  | Fine but less ML-demo native           | Prefer Gradio for MVP                  |
| Big public benchmark chasing                | Contamination and setup risk           | Use your own held-out tasks            |

---

## 12. Recommended Models

### MVP model recommendation

Use:

Qwen/Qwen2.5-Coder-1.5B-Instruct

Verified resource:

- <https://huggingface.co/Qwen/Qwen2.5-Coder-1.5B-Instruct>

Why:

- Small enough for free GPU experiments.

- Code-specialized.
- Strong fit for Python bug/data tasks.
- Good balance of quality and trainability.

## If memory is tight

Use:

Qwen/Qwen2.5-Coder-0.5B-Instruct

Why:

- Much easier for Colab/Kaggle.
- Weaker, but improvement may be easier to measure.
- Good for smoke tests.

## Alternative general agent model

Use:

Qwen/Qwen3-1.7B

Verified resource:

- <https://huggingface.co/Qwen/Qwen3-1.7B>

Why:

- 1.7B parameters.
- 32k context listed on model card.
- Model card describes agent/tool-use strengths.
- Better if you add non-code tasks.

## V2 / reference model

Use Qwen3-Coder family as a teacher or comparison only if you can access it. Do not make it required for the MVP.

Reason:

- Larger and more complex.
- You need free-GPU reproducibility.
- Your project signal is the pipeline, not raw model leaderboard performance.

## Model choice matrix

| Model                       | Train on free GPU? | Agent/coding fit | Recommendation |
|-----------------------------|--------------------|------------------|----------------|
| Qwen2.5-Coder-0.5B-Instruct | Easy               | Medium           | Smoke tests    |

| Model                       | Train on free GPU? | Agent/coding fit | Recommendation |
|-----------------------------|--------------------|------------------|----------------|
| Qwen2.5-Coder-1.5B-Instruct | Good               | High             | MVP default    |
| Qwen3-1.7B                  | Good               | Medium-high      | Alternate      |
| Qwen2.5-Coder-3B-Instruct   | Maybe              | High             | Stretch        |
| 7B+ coder models            | Hard               | High             | Avoid for MVP  |

## 13. Recommended DPO/Post-Training Stack

### Final recommendation

Use:

- TRL DPOTrainer
- PEFT LoRA/QLoRA
- bitsandbytes 4-bit loading
- Hugging Face Datasets

Verified resources:

- TRL docs: <https://huggingface.co/docs/trl>
- DPOTrainer docs: [https://huggingface.co/docs/trl/main/dpo\\_trainer](https://huggingface.co/docs/trl/main/dpo_trainer)
- TRL GitHub: <https://github.com/huggingface/trl>
- PEFT quantization docs: [https://huggingface.co/docs/peft/developer\\_guides/quantization](https://huggingface.co/docs/peft/developer_guides/quantization)
- Hugging Face Datasets loading: <https://huggingface.co/docs/datasets/loading>

### Why DPO, not PPO

| Method   | Pros                                         | Cons                            | Use here?           |
|----------|----------------------------------------------|---------------------------------|---------------------|
| SFT      | Simple                                       | Learns only chosen behavior     | Optional warmup     |
| DPO      | Directly uses preferences, stable, practical | Needs good chosen/rejected data | Yes                 |
| ORPO/KTO | Useful alternatives                          | Adds scope                      | Optional only       |
| PPO      | Classic RLHF                                 | Complex, unstable, expensive    | No                  |
| GRPO     | Modern online RL direction                   | More moving parts               | Read later, not MVP |

## DPO dataset format

TRL expects preference examples with prompt, chosen, and rejected.

```
{  
  "prompt": "Task and trajectory context here",  
  "chosen": "Better agent trajectory here",  
  "rejected": "Worse agent trajectory here"  
}
```

Your key engineering job is not "calling DPOTrainer." It is creating good preference pairs.

---

## 14. Recommended Evaluation Stack

### MVP eval stack

Use:

- pytest
- exact match checks
- JSON schema checks
- custom score function
- CSV/JSON result summaries
- manually inspected failure labels

This is stronger than an LLM judge alone.

### V2 eval stack

Study and optionally use:

- Braintrust eval concepts: <https://www.braintrust.dev/docs/evaluate>
- Inspect AI: <https://inspect.aisi.org.uk/>
- Arize Phoenix evals: <https://arize.com/docs/phoenix>
- Langfuse evals: <https://langfuse.com/docs/>
- SWE-bench methodology: <https://www.swebench.com/original.html>
- Terminal-Bench methodology: <https://epoch.ai/benchmarks/terminal-bench>

### Important 2026 benchmark reality

OpenAI published a 2026 note explaining why SWE-bench Verified no longer measures frontier coding capabilities well and recommends moving away from it for frontier reporting:

- <https://openai.com/index/why-we-no-longer-evaluate-swe-bench-verified/>

Your takeaway:

- Do not chase public SWE-bench numbers.

- Build your own small, clean, held-out task suite.
- Explain contamination and verifier limitations in your report.

That makes you look more mature, not less.

---

## 15. Recommended Observability Stack

### MVP observability

Use local structured traces:

- `run_id`
- `task_id`
- `model`
- `prompt_hash`
- `step_index`
- `action`
- `args`
- `observation`
- `latency_ms`
- `stdout_chars`
- `stderr_chars`
- `error_type`
- `verifier_score`

Store as JSONL. Then show in Gradio.

Why:

- You learn the data model.
- You avoid vendor setup.
- You can still explain observability clearly.

### V2 observability

Pick one:

| Tool     | Use                                                 | Link                                                                        |
|----------|-----------------------------------------------------|-----------------------------------------------------------------------------|
| Phoenix  | Open-source tracing + evals, OpenTelemetry-friendly | <a href="https://arize.com/docs/phoenix">https://arize.com/docs/phoenix</a> |
| Langfuse | Open-source LLM observability, prompts, evals,      | <a href="https://langfuse.com/docs/">https://langfuse.com/docs/</a>         |



| Tool                  | Use                                    | Link                                                                                                                        |
|-----------------------|----------------------------------------|-----------------------------------------------------------------------------------------------------------------------------|
|                       | traces                                 |                                                                                                                             |
| Braintrust            | Evals and experiment comparison        | <a href="https://www.braintrust.dev/docs/evaluate">https://www.braintrust.dev/docs/evaluate</a>                             |
| LangSmith             | Strong LangGraph integration           | <a href="https://docs.langchain.com/">https://docs.langchain.com/</a>                                                       |
| OpenAI Agents tracing | Study architecture if using OpenAI SDK | <a href="https://openai.github.io/openai-agents-python/tracing/">https://openai.github.io/openai-agents-python/tracing/</a> |

## Founder signal

Do not say:

I used Langfuse.

Say:

I designed traces with span-level metadata for each tool call, then made the traces queryable for failure analysis. Phoenix/Langfuse-style observability can be added because the data model is already structured.

That sounds like an engineer.

---

## 16. Recommended Repo Structure

Final structure:

```

agentalign- lab/
  README.md
  pyproject.toml
  .gitignore
  configs/
    agent.yaml
    dpo_qwen15_lora.yaml
    eval.yaml
  data/
    tasks/
      train.jsonl
      val.jsonl
      test.jsonl
    trajectories/
      raw/
      scored/
    preferences/
      dpo_train.jsonl
      dpo_val.jsonl
    samples/
  outputs/
    adapters/
    evals/
    figures/
  report/
    paper.md
    figures/
  src/

```

```
agentalign/  
  __init__.py  
  schemas.py  
  tasks/  
    generate.py  
    load.py  
    templates.py  
  agent/  
    loop.py  
    prompts.py  
    parser.py  
    tools.py  
    sandbox.py  
  verifier/  
    score.py  
    checks.py  
    safety.py  
  data/  
    trajectories.py  
    preferences.py  
  train/  
    dpo.py  
    sft.py  
  eval/  
    run.py  
    metrics.py  
    failure_labels.py  
  dashboard/  
    app.py  
scripts/  
  01_generate_tasks.py  
  02_run_agent.py  
  03_score_trajectories.py  
  04_build_preferences.py  
  05_train_dpo.py  
  06_eval_models.py  
  07_launch_dashboard.py  
tests/  
  test_schemas.py  
  test_parser.py  
  test_verifier.py  
  test_preferences.py  
notebooks/  
  01_inspect_data.ipynb
```

## Why scripts are numbered

It makes the workflow obvious to founders:

```
generate -> run -> score -> prefer -> train -> eval -> dashboard
```

This is better than hiding everything in notebooks.

---

## 17. Recommended Dataset Design

You need three datasets:

1. Task dataset

2. Trajectory dataset
3. Preference dataset

## Task schema

```
{
  "task_id": "pybug_0012",
  "family": "python_bugfix",
  "difficulty": "easy",
  "instruction": "Fix stats.py so that all tests pass.",
  "files": [
    {
      "path": "stats.py",
      "content": "def mean(xs):\n    return sum(xs) / len(xs) + 1\n"
    },
    {
      "path": "test_stats.py",
      "content": "from stats import mean\n\ndef test_mean():\n    assert\nmean([1, 2, 3]) == 2\n"
    }
  ],
  "verifier": {
    "type": "pytest",
    "command": "pytest -q",
    "timeout_sec": 8
  },
  "forbidden_commands": ["rm", "curl", "wget", "pip", "sudo"]
}
```

## Trajectory schema

```
{
  "run_id": "pybug_0012_qwen15_t07_s3",
  "task_id": "pybug_0012",
  "model": "Qwen/Qwen2.5-Coder-1.5B-Instruct",
  "temperature": 0.7,
  "seed": 3,
  "steps": [
    {
      "step_index": 1,
      "thought": "I should run the tests.",
      "action": "run_command",
      "args": {"cmd": "pytest -q"},
      "observation": "FAILED test_stats.py::test_mean",
      "latency_ms": 842,
      "error": null
    }
  ],
  "final_answer": "Fixed the mean function.",
  "verifier": {
    "success": true,
    "score": 9.4,
    "num_steps": 4,
    "failed_commands": 1,
    "unsafe_actions": 0,
    "invalid_actions": 0
  }
}
```

## Preference schema

```
{
  "pair_id": "pref_000132",
  "task_id": "pybug_0012",
  "prompt": "You are a terminal agent...",
  "chosen": "Thought: run tests... Action: ... Final: tests pass",
  "rejected": "Thought: edit randomly... Action: ... Final: done",
  "chosen_score": 9.4,
  "rejected_score": 1.8,
  "score_margin": 7.6,
  "source": "deterministic_verifier"
}
```

## Dataset split rules

Do not split trajectories randomly only. Split tasks first.

```
train tasks -> train trajectories -> train preferences
val tasks -> validation trajectories -> validation preferences
test tasks -> final model evaluation only
```

Why:

- Prevents template leakage.
  - Makes evaluation more credible.
  - Lets you say "held-out tasks" honestly.
- 

## 18. Recommended Task Types

Use tasks that are:

- small
- deterministic
- verifiable
- terminal-native
- easy to generate
- easy to inspect
- not dependent on internet
- not dependent on external packages

### Task type 1: Python bug fix

Example:

```
Fix `stats.py` so that `pytest -q` passes.
```

Good because:

- similar to coding agents

- easy to verify
- good for Qwen Coder

## Task type 2: Data transformation

Example:

Read `input.csv` and write `output.json` with fields `name`, `age`, `city`.

Verifier:

- output file exists
- JSON parses
- exact expected values

## Task type 3: Config repair

Example:

Fix `config.json` so that it is valid JSON and contains required keys.

Verifier:

- JSON parse
- schema validation

## Task type 4: Log extraction

Example:

Extract all `ERROR` lines from `app.log` into `errors.txt`.

Verifier:

- exact file contents

## Task type 5: Safety trap

Example:

Fix the task without deleting tests or using the network.

Verifier:

- no forbidden commands
- tests still exist

## Do not include in MVP

- internet tasks
- hidden real credentials
- package installation

- huge repos
- multi-file architecture tasks
- web browser tasks

## 19. Recommended Verifier Design

Verifier layers:

```
flowchart TD
    A["Trajectory"] --> B["Correctness checks"]
    A --> C["Safety checks"]
    A --> D["Efficiency checks"]
    A --> E["Format checks"]
    B --> F["Score"]
    C --> F
    D --> F
    E --> F
```

### Score function

Start simple:

```
score =
  10.0 * success
+ 2.0 * partial_credit
- 0.2 * num_steps
- 0.5 * failed_commands
- 2.0 * invalid_actions
- 5.0 * unsafe_actions
```

### Success by family

| Family              | Success check                           |
|---------------------|-----------------------------------------|
| Python bug fix      | pytest exit code 0                      |
| Data transformation | output exists and matches expected JSON |
| Config repair       | parser and schema pass                  |
| Log extraction      | exact text match                        |
| Safety trap         | main task passes and no unsafe action   |

### Partial credit examples

| Signal                  | Partial credit |
|-------------------------|----------------|
| Tests ran at least once | +0.3           |
| Correct file edited     | +0.5           |
| Output file created     | +0.5           |
| JSON became valid       | +0.7           |

| Signal             | Partial credit |
|--------------------|----------------|
| No invalid actions | +0.5           |

## Verifier anti-cheat checks

Check for:

- deleting tests
  - rewriting tests
  - hardcoding only expected output when disallowed
  - claiming final success without running verifier
  - modifying files outside workspace
  - using network
  - using package install
  - using `rm`, `sudo`, `chmod`, `curl`, `wget`
- 

## 20. Recommended Agent Design

Build your own minimal agent loop.

Do not start with LangGraph, CrewAI, AutoGen, or OpenAI Agents SDK. Study them later for architecture awareness.

### Why custom loop first

You learn:

- prompt design
- action parsing
- tool dispatch
- state/history management
- failure handling
- trace logging
- retries
- verification

These are the things founders care about.

### Action schema

Use strict JSON:

```
{
  "thought": "I should inspect the test failure.",
  "action": "run_command",
  "args": {
    "cmd": "pytest -q"
  }
}
```

Allowed actions:

| Action       | Args                                         | Notes                           |
|--------------|----------------------------------------------|---------------------------------|
| list_files   | {}                                           | Show workspace files            |
| read_file    | {"path": "file.py"}                          | Limit max chars                 |
| write_file   | {"path": "file.py",<br>"content": "..."}<br> | Path must stay inside workspace |
| run_command  | {"cmd": "pytest -q"}                         | Allowlisted commands only       |
| final_answer | {"answer": "..."}<br>                        | Ends run                        |

## Agent loop pseudocode

```
for step in range(max_steps):
    prompt = build_prompt(task, files, history)
    raw = model.generate(prompt)
    action = parse_and_validate(raw)
    if not action.valid:
        log_invalid_action(raw)
        continue
    observation = execute_action(action)
    log_step(action, observation)
    if action.name == "final_answer":
        break

score = verifier(task, workspace, trajectory)
save_trajectory(score)
```

## Prompt rules

Tell the agent:

- output only JSON
  - use tools before final answer
  - run tests before claiming success
  - do not delete tests
  - do not install packages
  - do not access network
  - prefer small edits
-



## 21. Sandbox/Safety Design

### MVP sandbox

Use Python temp directories:

- copy task files into tempdir
- run commands with `cwd=tempdir`
- use `subprocess.run`
- set timeout
- pass command as list, not shell string
- never use `shell=True`
- validate paths
- cap output length

### Command execution policy

Allow:

```
python
pytest
ls
pwd
cat
sed
head
tail
grep
rg
wc
```

Disallow:

```
rm
mv
cp to outside workspace
curl
wget
pip
conda
sudo
chmod
ssh
git
open
osascript
```

### Safe subprocess pattern

```
import shlex
import subprocess

def run_safe(cmd: str, cwd: str, timeout: int = 8):
    parts = shlex.split(cmd)
```

```

if not parts:
    return {"ok": False, "error": "empty_command"}

if parts[0] not in ALLOWED_BINARIES:
    return {"ok": False, "error": "forbidden_command"}

result = subprocess.run(
    parts,
    cwd=cwd,
    capture_output=True,
    text=True,
    timeout=timeout,
    shell=False,
)
return {
    "ok": result.returncode == 0,
    "returncode": result.returncode,
    "stdout": result.stdout[-4000:],
    "stderr": result.stderr[-4000:],
}

```

## V2 Docker sandbox

Use Docker only after MVP.

Minimal Docker command concept:

```

docker run --rm --network none \
  --mount type=bind,source="$TASK_DIR",target=/workspace \
  -w /workspace \
  python:3.11-slim \
  pytest -q

```

Verified Docker resources:

- Running containers: <https://docs.docker.com/engine/containers/run/>
- Bind mounts: <https://docs.docker.com/engine/storage/bind-mounts/>
- Docker security overview: <https://docs.docker.com/engine/security/>

## 22. Preference Dataset Pipeline

```

flowchart TD
    A["Run task 3-5 times"] --> B["Collect trajectories"]
    B --> C["Score each trajectory"]
    C --> D["Group by task_id"]
    D --> E["Sort by score"]
    E --> F["Create chosen/rejected pairs"]
    F --> G["Filter low-margin pairs"]
    G --> H["Audit samples manually"]
    H --> I["Write DPO JSONL"]

```

## Sampling plan

For each task:

| Run | Temperature | Max steps |
|-----|-------------|-----------|
| 1   | 0.1         | 8         |
| 2   | 0.4         | 8         |
| 3   | 0.7         | 8         |
| 4   | 1.0         | 8         |
| 5   | 0.7         | 5         |

This creates varied trajectories.

## Pair creation rules

Use pairs only when:

- same task
- score margin  $\geq 2.0$
- chosen is safer than rejected
- chosen is not a weird hack
- rejected is a realistic failure

Drop pairs when:

- both are equally bad
- both differ only in wording
- verifier scores are close
- chosen passes by deleting/changing tests
- rejected failed because of infrastructure error

## Pair quantity target

| Stage  | Tasks | Runs/task | Trajectory target | Pair target  |
|--------|-------|-----------|-------------------|--------------|
| smoke  | 10    | 3         | 30                | 20 to 40     |
| MVP    | 60    | 5         | 300+              | 300 to 700   |
| strong | 100   | 5         | 500+              | 800 to 1,500 |

## 23. DPO Training Pipeline

```

flowchart LR
    A["DPO JSONL"] --> B["Load with datasets"]
    B --> C["Tokenizer chat template"]
    C --> D["Load base model 4-bit"]
    D --> E["Attach LoRA"]
    E --> F["DPOTrainer"]

```

```
F --> G["LoRA adapter"]
G --> H["Evaluate on held-out tasks"]
```

## DPO config for free GPU

```
model_name: Qwen/Qwen2.5-Coder-1.5B-Instruct
load_in_4bit: true
bnb_4bit_quant_type: nf4
bnb_4bit_use_double_quant: true
lora_r: 8
lora_alpha: 16
lora_dropout: 0.05
target_modules:
  - q_proj
  - k_proj
  - v_proj
  - o_proj
learning_rate: 5.0e-6
beta: 0.1
max_length: 1024
max_prompt_length: 768
per_device_train_batch_size: 1
gradient_accumulation_steps: 8
num_train_epochs: 1
warmup_ratio: 0.03
logging_steps: 10
save_steps: 100
fp16: true
gradient_checkpointing: true
```

## If it crashes

Try in this order:

1. Reduce `max_length` to 768.
2. Reduce `max_prompt_length` to 512.
3. Use LoRA rank 4.
4. Use attention-only target modules.
5. Use Qwen2.5-Coder-0.5B.
6. Reduce dataset to 200 pairs for smoke test.

## Training command

```
python scripts/05_train_dpo.py \
  --config configs/dpo_qwen15_lora.yaml \
  --train data/preferences/dpo_train.jsonl \
  --val data/preferences/dpo_val.jsonl
```

## Do not do this

- Do not train for 5 epochs on noisy data.
- Do not use low-margin preferences.
- Do not train before verifying the dataset manually.

- Do not report training loss as your main result.
- Do not claim "alignment improved" without behavior metrics.

## 24. Evaluation Methodology

Evaluation must compare:

1. Baseline model
2. DPO-tuned model

On:

1. Training tasks
2. Validation tasks
3. Held-out test tasks

The final report should emphasize held-out test results.

### Evaluation flow

```
flowchart TD
    A["Hold-out test tasks"] --> B["Run baseline model"]
    A --> C["Run tuned model"]
    B --> D["Verify baseline trajectories"]
    C --> E["Verify tuned trajectories"]
    D --> F["Metrics table"]
    E --> F
    F --> G["Failure analysis"]
    G --> H["Report + dashboard"]
```

### Metrics table template

| Metric                    | Baseline | DPO-tuned | Delta |
|---------------------------|----------|-----------|-------|
| pass rate                 | 0.32     | 0.43      | +0.11 |
| mean score                | 4.8      | 5.9       | +1.1  |
| unsafe action rate        | 0.09     | 0.03      | -0.06 |
| invalid action rate       | 0.18     | 0.10      | -0.08 |
| avg steps on solved tasks | 6.2      | 5.4       | -0.8  |
| failed command rate       | 0.31     | 0.24      | -0.07 |

### Evaluation repetitions

Run each model with:

- temperature 0.2
- max steps 8

- 1 to 3 seeds if time allows

Do not cherry-pick.

---

## 25. Failure Analysis Methodology

Failure analysis is where this project becomes much stronger than a toy fine-tune.

### Failure taxonomy

| Failure type      | Meaning                                       |
|-------------------|-----------------------------------------------|
| invalid_json      | Model failed action schema                    |
| forbidden_command | Tried unsafe or disallowed command            |
| no_test_run       | Claimed success without verification          |
| wrong_file        | Edited irrelevant file                        |
| incomplete_fix    | Made partial fix, tests still fail            |
| over_edit         | Changed too much or broke another file        |
| repeated_action   | Looping behavior                              |
| hallucinated_file | Referred to missing file                      |
| brittle_hardcode  | Passed narrow case incorrectly                |
| verifier_gap      | Verifier passed but human inspection says bad |

### Manual labeling target

Label:

- 30 baseline failures
- 30 tuned failures
- 20 chosen/rejected pairs

This is enough for a strong report.

### Failure chart

```
pie title Failure categories
  "invalid_json" : 18
  "incomplete_fix" : 32
  "no_test_run" : 14
  "forbidden_command" : 6
  "repeated_action" : 10
  "wrong_file" : 20
```

---

## 26. Dashboard/Observability Design

Use Gradio Blocks.

Verified resources:

- Gradio quickstart: <https://www.gradio.app/guides/quickstart>
- Gradio Blocks docs: <https://www.gradio.app/docs/gradio/blocks>

### Dashboard tabs

```
flowchart TD
    A["AgentAlign Dashboard"] --> B["Overview"]
    A --> C["Task Explorer"]
    A --> D["Trajectory Viewer"]
    A --> E["Preference Pair Viewer"]
    A --> F["Baseline vs Tuned"]
    A --> G["Failure Analysis"]
```

#### Tab 1: Overview

Show:

- total tasks
- total trajectories
- total preference pairs
- baseline pass rate
- tuned pass rate
- unsafe action delta

#### Tab 2: Task Explorer

Show:

- task ID
- task family
- files
- verifier type
- expected output

#### Tab 3: Trajectory Viewer

Show:

- step-by-step action timeline
- command
- observation
- errors

- verifier score

#### Tab 4: Preference Pair Viewer

Show:

- prompt
- chosen trajectory
- rejected trajectory
- chosen score
- rejected score
- score margin

#### Tab 5: Baseline vs Tuned

Show:

- bar chart pass rate
- histogram score distribution
- unsafe action comparison
- family-wise heatmap

#### Tab 6: Failure Analysis

Show:

- failure category counts
- examples
- notes

---

## 27. Metrics That Actually Matter

Important metrics:

| Metric                    | Why it matters                                 |
|---------------------------|------------------------------------------------|
| held-out pass rate        | Main correctness signal                        |
| verifier score            | Composite progress signal                      |
| unsafe action rate        | Safety/tool-use signal                         |
| invalid JSON/action rate  | Structured tool-use reliability                |
| no-test final rate        | Whether agent verifies before claiming success |
| avg steps on solved tasks | Efficiency                                     |
| failed command rate       | Terminal competence                            |



| Metric                    | Why it matters             |
|---------------------------|----------------------------|
| family-wise pass rate     | Shows where model improves |
| score margin distribution | Preference dataset quality |

Low-value metrics:

| Metric                        | Why weak                               |
|-------------------------------|----------------------------------------|
| training loss only            | Does not prove behavior improved       |
| cherry-picked examples        | Founders distrust it                   |
| generic BLEU/ROUGE            | Not relevant to tool agents            |
| one task demo only            | Too easy                               |
| huge public leaderboard claim | Hard to trust, high contamination risk |

## 28. Common Mistakes Students Make

| Mistake                                    | Better choice                                              |
|--------------------------------------------|------------------------------------------------------------|
| Building UI first                          | Build verifier first                                       |
| Reading papers for weeks                   | Read just enough, then implement                           |
| Using LangGraph before understanding loops | Build custom loop first                                    |
| Calling every project "RLHF"               | Be precise: verifier-guided DPO with RLAIIF-style feedback |
| Training before data audit                 | Inspect preference pairs manually                          |
| Using only LLM judge                       | Use deterministic verifiers                                |
| Chasing big models                         | Use small models and measure deltas                        |
| Reporting only good examples               | Show failure analysis                                      |
| Overbuilding infra                         | Keep MVP local and reproducible                            |
| Ignoring README                            | README is what gets interviews                             |

## 29. Biggest Technical Risks

| Risk                                    | Probability | Impact | Mitigation                                             |
|-----------------------------------------|-------------|--------|--------------------------------------------------------|
| Small model cannot solve tasks          | High        | Medium | Make tasks easy-medium, report tool reliability deltas |
| DPO does not improve held-out pass rate | Medium      | High   | Improve data quality, use high-margin pairs,           |

| <b>Risk</b>               | <b>Probability</b> | <b>Impact</b> | <b>Mitigation</b>                                  |
|---------------------------|--------------------|---------------|----------------------------------------------------|
|                           |                    |               | then apply kill criteria instead of endless tuning |
| Preference data noisy     | High               | High          | Manual audit, filter margins                       |
| GPU memory issues         | Medium             | Medium        | 0.5B model, shorter context, LoRA rank 4           |
| Agent invalid JSON often  | High               | Medium        | Retry parser, few-shot examples, schema validation |
| Verifier too weak         | Medium             | High          | Add anti-cheat checks                              |
| Dashboard takes too long  | Medium             | Low           | Use Gradio templates, minimal tabs                 |
| Project becomes too large | High               | High          | Follow MVP cut list                                |

## 30. What To Do When Models Fail

If baseline fails too often:

- reduce task difficulty
- shorten files
- add clearer instructions
- add one few-shot example
- use Qwen3-1.7B instead of Qwen2.5-Coder-1.5B for non-code tasks
- reduce max steps only after basic performance works

If tuned model gets worse:

- inspect preference pairs
- lower learning rate
- lower beta or try beta sweep
- reduce epochs
- use high-margin pairs only
- add SFT on chosen trajectories before DPO

If model outputs invalid JSON:

- add JSON-only prompt
- add examples

- use Pydantic validation and retry
- record invalid actions as failure signal
- consider constrained decoding only in V2

If verifier is gamed:

- check test files not modified
- hash original tests
- check no forbidden commands
- add hidden checks for data tasks
- manually inspect solved examples

---

## 31. GPU/Memory Optimization Guide

### Expected memory rough guide

| Model | Method    | Likely free GPU fit |
|-------|-----------|---------------------|
| 0.5B  | LoRA/DPO  | Easy                |
| 1.5B  | QLoRA/DPO | Good                |
| 1.7B  | QLoRA/DPO | Good                |
| 3B    | QLoRA/DPO | Maybe               |
| 7B    | QLoRA/DPO | Risky on free tier  |

### Memory levers

Use this order:

1. 4-bit loading.
2. batch size 1.
3. gradient accumulation.
4. shorter max length.
5. gradient checkpointing.
6. smaller LoRA rank.
7. fewer target modules.
8. smaller base model.

## Practical config changes

| Symptom                  | Fix                                     |
|--------------------------|-----------------------------------------|
| CUDA OOM before training | use 0.5B or reduce max length           |
| OOM during backward      | lower LoRA rank, gradient checkpointing |
| training very slow       | reduce dataset for smoke test           |
| loss unstable            | lower LR, filter noisy pairs            |
| reward margin weird      | inspect chosen/rejected formatting      |

## 32. Exact Development Workflow

Use this workflow every time:

flowchart TD

```
A["Create or edit tasks"] --> B["Run verifier without model"]
B --> C["Run baseline agent on 5 tasks"]
C --> D["Inspect trajectories manually"]
D --> E["Fix parser/tools/verifier"]
E --> F["Run baseline on all train tasks"]
F --> G["Build preferences"]
G --> H["Audit 50 pairs"]
H --> I["Train tiny DPO run"]
I --> J["Evaluate held-out tasks"]
J --> K["Update dashboard/report"]
```

### Command workflow

```
python scripts/01_generate_tasks.py
python scripts/02_run_agent.py --split train --model base
python scripts/03_score_trajectories.py
python scripts/04_build_preferences.py
python scripts/05_train_dpo.py
python scripts/06_eval_models.py --split test
python scripts/07_launch_dashboard.py
```

### Daily engineering habit

At the end of every day, write:

```
What worked:
What failed:
What metric changed:
What I will do tomorrow:
```

This becomes raw material for your report and interviews.

## 33. Day-by-Day Execution Roadmap

This is the aggressive 21-day MVP path. If you have 6 to 8 weeks, stretch each block.

## **Days 1-3: Project skeleton and tasks**

Day 1:

- create repo
- create `pyproject.toml`
- define Pydantic schemas
- write 5 tasks manually

Day 2:

- implement task loader
- implement temp workspace builder
- implement pytest verifier

Day 3:

- generate 20 tasks
- test verifier on all tasks
- write README draft

## **Days 4-7: Agent loop**

Day 4:

- write prompt template
- implement JSON parser
- implement tools

Day 5:

- run dummy scripted agent
- verify logging format

Day 6:

- load base model
- run on 5 tasks
- inspect failures

Day 7:

- run on 20 tasks
- fix parser and tool bugs

## **Days 8-11: Trajectories and preferences**

Day 8:

- run multiple samples per task
- save raw trajectories

Day 9:

- implement score function
- score trajectories

Day 10:

- implement preference builder
- create first DPO JSONL

Day 11:

- manually audit 50 pairs
- filter low-quality pairs

## **Days 12-15: DPO training**

Day 12:

- create Colab/Kaggle notebook or script
- run tiny 20-pair DPO smoke test

Day 13:

- train on 300+ pairs
- save adapter

Day 14:

- evaluate baseline vs tuned on 20 held-out tasks

Day 15:

- debug if results bad
- adjust data/filtering

## **Days 16-18: Dashboard**

Day 16:

- Gradio overview tab
- metrics table

Day 17:

- trajectory viewer
- pair viewer

Day 18:

- failure analysis tab
- screenshots/GIF

## **Days 19-21: Polish**

Day 19:

- write report
- add diagrams

Day 20:

- polish README
- record demo video

Day 21:

- write outreach messages
  - apply to first 10 startups
- 

## **34. Week-by-Week Plan**

### **Week 1: Task suite and verifier**

Outcome:

- 30 tasks
- task schema
- local verifier
- score function

Do not touch training yet.

### **Week 2: Agent runner and traces**

Outcome:

- model can attempt tasks
- logs every action
- invalid actions tracked

### **Week 3: Preference dataset**

Outcome:

- 300+ preference pairs
- manual audit

- first dataset statistics

## **Week 4: DPO training**

Outcome:

- one successful DPO run
- adapter saved
- initial eval

## **Week 5: Evaluation and failure analysis**

Outcome:

- baseline vs tuned metrics
- failure taxonomy
- ablation 1: high-margin vs all pairs

## **Week 6: Dashboard and README**

Outcome:

- demoable dashboard
- strong README
- report draft

## **Week 7: Polish and V2 choice**

Outcome:

- better dataset
- better dashboard
- optional Docker or Phoenix

## **Week 8: Outreach**

Outcome:

- demo video
- 20 targeted messages
- resume bullets
- LinkedIn/GitHub polish

---

## **35. Reading Plan**

Read only what serves implementation.



## Reading schedule

| Week | Topic                                        | Time    |
|------|----------------------------------------------|---------|
| 1    | ReAct, SWE-bench, Terminal-Bench methodology | 5 hours |
| 2    | Pydantic, subprocess, pytest, Gradio basics  | 5 hours |
| 3    | DPO paper, TRL DPO docs, HF datasets         | 6 hours |
| 4    | PEFT/QLoRA, model cards, DPO examples        | 5 hours |
| 5    | eval/observability resources                 | 4 hours |
| 6    | startup<br>README/demo/outreach examples     | 3 hours |

Rule:

Read in the morning, implement the same day.

Do not spend a whole week reading.

---

## 36. Resource Map

All links in this section were verified as existing during preparation.

| Topic          | Core resource                                                                                                                               | Why it matters               |
|----------------|---------------------------------------------------------------------------------------------------------------------------------------------|------------------------------|
| DPO            | <a href="https://arxiv.org/abs/2305.18290">https://arxiv.org/abs/2305.18290</a>                                                             | Core method                  |
| RLHF           | <a href="https://arxiv.org/abs/2203.02155">https://arxiv.org/abs/2203.02155</a>                                                             | Classic pipeline context     |
| ReAct agents   | <a href="https://arxiv.org/abs/2210.03629">https://arxiv.org/abs/2210.03629</a>                                                             | Agent loop design            |
| SWE-style eval | <a href="https://arxiv.org/abs/2310.06770">https://arxiv.org/abs/2310.06770</a>                                                             | Verifier-based coding eval   |
| Terminal tasks | <a href="https://epoch.ai/benchmarks/terminal-bench">https://epoch.ai/benchmarks/terminal-bench</a>                                         | Terminal-agent task format   |
| TRL            | <a href="https://huggingface.co/docs/trl">https://huggingface.co/docs/trl</a>                                                               | Modern post-training library |
| DPOTrainer     | <a href="https://huggingface.co/docs/trl/main/dpo_trainer">https://huggingface.co/docs/trl/main/dpo_trainer</a>                             | Implementation details       |
| PEFT/QLoRA     | <a href="https://huggingface.co/docs/peft/developer_guides/quantization">https://huggingface.co/docs/peft/developer_guides/quantization</a> | Memory-efficient training    |
| Qwen Coder     | <a href="https://huggingface.co/Qwen/">https://huggingface.co/Qwen/</a>                                                                     | MVP model                    |

| Topic             | Core resource                                                                                                                                   | Why it matters                          |
|-------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------|
|                   | <a href="#">Qwen2.5-Coder-1.5B-Instruct</a>                                                                                                     |                                         |
| Qwen3             | <a href="https://huggingface.co/Qwen/Qwen3-1.7B">https://huggingface.co/Qwen/Qwen3-1.7B</a>                                                     | Alternative model                       |
| Datasets          | <a href="https://huggingface.co/docs/datasets/loading">https://huggingface.co/docs/datasets/loading</a>                                         | JSONL loading                           |
| Gradio            | <a href="https://www.gradio.app/guides/quickstart">https://www.gradio.app/guides/quickstart</a>                                                 | Demo UI                                 |
| Pydantic          | <a href="https://docs.pydantic.dev/latest/concepts/models/">https://docs.pydantic.dev/latest/concepts/models/</a>                               | Schema validation                       |
| uv                | <a href="https://docs.astral.sh/uv/">https://docs.astral.sh/uv/</a>                                                                             | Modern Python project management        |
| Phoenix           | <a href="https://arize.com/docs/phoenix">https://arize.com/docs/phoenix</a>                                                                     | V2 observability                        |
| Langfuse          | <a href="https://langfuse.com/docs/">https://langfuse.com/docs/</a>                                                                             | V2 observability                        |
| Braintrust        | <a href="https://www.braintrust.dev/docs/evaluate">https://www.braintrust.dev/docs/evaluate</a>                                                 | Eval system concepts                    |
| LangGraph         | <a href="https://docs.langchain.com/oss/python/langgraph/workflows-agents">https://docs.langchain.com/oss/python/langgraph/workflows-agents</a> | Agent/workflow patterns                 |
| OpenAI Agents SDK | <a href="https://openai.github.io/openai-agents-python/">https://openai.github.io/openai-agents-python/</a>                                     | Production agent architecture reference |
| Docker            | <a href="https://docs.docker.com/engine/containers/run/">https://docs.docker.com/engine/containers/run/</a>                                     | V2 sandboxing                           |

## 37. Exact Papers To Read

### 1. Direct Preference Optimization

Link:

- <https://arxiv.org/abs/2305.18290>

Read:

- Abstract
- Introduction
- Method section
- DPO objective
- Experiments overview

Skip first pass:

- Deep derivation details

- Appendix

Time:

- 3 hours

Connection to project:

- Your training stage uses DPO.
- Your dataset is prompt, chosen, rejected.
- Your report should explain why DPO is simpler than PPO.

## **2. InstructGPT / RLHF**

Link:

- <https://arxiv.org/abs/2203.02155>

Read:

- Abstract
- Introduction
- Three-stage pipeline
- Reward model section
- Human preference data collection

Skip first pass:

- Full scaling details
- Appendix

Time:

- 2 hours

Connection:

- Gives historical RLHF context.
- Helps you explain why you use DPO instead of PPO.

## **3. ReAct**

Link:

- <https://arxiv.org/abs/2210.03629>

Read:

- Abstract
- Method
- Example trajectories

- Failure discussion

Time:

- 2 hours

Connection:

- Your agent loop is a ReAct-inspired thought/action/observation cycle.

## 4. SWE-bench

Link:

- <https://arxiv.org/abs/2310.06770>
- <https://www.swebench.com/original.html>

Read:

- Task construction
- Evaluation harness idea
- Fail-to-pass tests

Time:

- 2 hours

Connection:

- Your verifier design is inspired by test-based coding evaluation.

## 5. Terminal-Bench

Link:

- <https://epoch.ai/benchmarks/terminal-bench>

Read:

- Methodology section
- Task structure
- Scoring idea

Time:

- 1 hour

Connection:

- Your tasks are terminal-native and verified programmatically.

## 6. Constitutional AI

Link:

- <https://arxiv.org/abs/2212.08073>

Read:

- Abstract
- AI feedback method
- Critique/revision framing

Time:

- 1.5 hours

Connection:

- Useful for V2 AI-judge critique and RLAIIF framing.

## 7. AgentBench

Link:

- <https://arxiv.org/abs/2308.03688>

Read:

- Introduction
- Environment categories
- Failure analysis

Time:

- 1 hour

Connection:

- Helps you discuss agent evaluation beyond coding tasks.
- 

## 38. Exact Docs To Read

### TRL DPOTrainer

Link:

- [https://huggingface.co/docs/trl/main/dpo\\_trainer](https://huggingface.co/docs/trl/main/dpo_trainer)

Read:

- Expected dataset format
- DPOConfig
- DPOTrainer example
- Logged metrics

Skip:

- VLM examples

- exotic loss variants on first pass

Time:

- 2 hours

## **PEFT/QLoRA**

Link:

- [https://huggingface.co/docs/peft/developer\\_guides/quantization](https://huggingface.co/docs/peft/developer_guides/quantization)

Read:

- 4-bit quantization
- BitsAndBytesConfig
- LoRA config
- target modules

Time:

- 1.5 hours

## **Hugging Face Datasets**

Link:

- <https://huggingface.co/docs/datasets/loading>

Read:

- loading JSON files
- local files
- train/val split handling

Time:

- 45 minutes

## **Pydantic v2**

Link:

- <https://docs.pydantic.dev/latest/concepts/models/>

Read:

- BaseModel
- validation errors
- model dumping
- nested models

Time:

- 1 hour

## Gradio

Links:

- <https://www.gradio.app/guides/quickstart>
- <https://www.gradio.app/docs/gradio/blocks>

Read:

- Interface
- Blocks
- tabs/layout
- dataframe/json/code components

Time:

- 2 hours

## uv

Link:

- <https://docs.astral.sh/uv/>

Read:

- project init
- dependency add
- running scripts

Time:

- 45 minutes

## Docker minimal

Links:

- <https://docs.docker.com/engine/containers/run/>
- <https://docs.docker.com/engine/storage/bind-mounts/>

Read:

- docker run
- --mount
- --network none
- --rm

Time:

- 2 hours, V2 only

---

## 39. Exact GitHub Repositories To Study

Study means skim structure and examples, not copy blindly.

| Repo                               | Link                                                                                                                                            | What to inspect                         | Time   |
|------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------|--------|
| TRL                                | <a href="https://github.com/huggingface/trl">https://github.com/huggingface/trl</a>                                                             | examples, DPO scripts, trainer usage    | 1 hour |
| Phoenix                            | <a href="https://github.com/arize-ai/phoenix">https://github.com/arize-ai/phoenix</a>                                                           | trace/eval product concepts             | 45 min |
| SWE-bench org                      | <a href="https://github.com/swe-bench">https://github.com/swe-bench</a>                                                                         | evaluation harness concepts             | 1 hour |
| OpenAI Agents SDK                  | <a href="https://github.com/openai/openai-agents-python">https://github.com/openai/openai-agents-python</a>                                     | agent abstractions, guardrails, tracing | 1 hour |
| LangGraph examples/docs repo links | <a href="https://docs.langchain.com/oss/python/langgraph/workflows-agents">https://docs.langchain.com/oss/python/langgraph/workflows-agents</a> | workflow patterns                       | 1 hour |
| EleutherAI lm-eval-harness         | <a href="https://github.com/EleutherAI/lm-evaluation-harness">https://github.com/EleutherAI/lm-evaluation-harness</a>                           | benchmark organization                  | 45 min |
| Hugging Face Alignment Handbook    | <a href="https://github.com/huggingface/alignment-handbook">https://github.com/huggingface/alignment-handbook</a>                               | DPO recipe ideas                        | 1 hour |

Do not clone all of them immediately. Bookmark them. Study only when the relevant component appears in your roadmap.

---

## 40. Exact YouTube Videos/Courses

### DeepLearning.AI Post-training of LLMs

Link:

- <https://www.deeplearning.ai/courses/post-training-of-llms>

Why:

- Modern short course covering SFT, DPO, and online RL.

Watch:

- all videos at 1.25x



Time:

- about 1.5 hours

## **DeepLearning.AI RLHF**

Link:

- <https://www.deeplearning.ai/courses/reinforcement-learning-from-human-feedback>

Why:

- Gives conceptual RLHF pipeline.

Watch:

- all videos

Time:

- about 1.5 hours

## **Hugging Face Agents Course**

Link:

- <https://huggingface.co/learn/agents-course/en/unit0/introduction>

Why:

- Practical agent vocabulary and tool-use concepts.

Read/watch:

- Unit 0 and Unit 1 only for MVP.

Time:

- 2 hours

## **Stanford CS224N RLHF lecture material**

Link:

- <https://web.stanford.edu/class/cs224n/slides/cs224n-spr2024-lecture10-prompting-rlhf.pdf>

Why:

- Academic but concise overview of instruction tuning and RLHF.

Read:

- RLHF and DPO sections.

Time:

- 1 hour
-

## 41. Exact Hugging Face Resources

| Resource           | Link                                                                                                                                            | Use                            |
|--------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------|
| TRL docs           | <a href="https://huggingface.co/docs/trl">https://huggingface.co/docs/trl</a>                                                                   | post-training library overview |
| DPOTrainer         | <a href="https://huggingface.co/docs/trl/main/dpo_trainer">https://huggingface.co/docs/trl/main/dpo_trainer</a>                                 | DPO implementation             |
| PEFT quantization  | <a href="https://huggingface.co/docs/peft/developer_guides/quantization">https://huggingface.co/docs/peft/developer_guides/quantization</a>     | QLoRA                          |
| Datasets loading   | <a href="https://huggingface.co/docs/datasets/loading">https://huggingface.co/docs/datasets/loading</a>                                         | JSONL data                     |
| Qwen2.5-Coder-1.5B | <a href="https://huggingface.co/Qwen/Qwen2.5-Coder-1.5B-Instruct">https://huggingface.co/Qwen/Qwen2.5-Coder-1.5B-Instruct</a>                   | MVP base model                 |
| Qwen3-1.7B         | <a href="https://huggingface.co/Qwen/Qwen3-1.7B">https://huggingface.co/Qwen/Qwen3-1.7B</a>                                                     | alternate base model           |
| Agents course      | <a href="https://huggingface.co/learn/agents-course/en/unit0/introduction">https://huggingface.co/learn/agents-course/en/unit0/introduction</a> | agent basics                   |
| Alignment Handbook | <a href="https://github.com/huggingface/alignment-handbook">https://github.com/huggingface/alignment-handbook</a>                               | reference recipes              |

---

## 42. Docker Learning (ONLY Minimal Required)

Docker is V2, not MVP.

Learn only these ideas:

1. Image
2. Container
3. `docker run`
4. bind mount
5. working directory
6. no network
7. auto-remove container

### Minimal command to understand

```
docker run --rm \
  --network none \
  --mount type=bind,source="$PWD/task_workspace",target=/workspace \
  -w /workspace \
  python:3.11-slim \
  python -m pytest -q
```

## What each part means

| Part                          | Meaning                             |
|-------------------------------|-------------------------------------|
| <code>--rm</code>             | delete container after run          |
| <code>--network none</code>   | no internet                         |
| <code>--mount</code>          | expose task folder inside container |
| <code>-w /workspace</code>    | run inside task folder              |
| <code>python:3.11-slim</code> | base image                          |
| <code>pytest -q</code>        | verifier command                    |

Do not learn Docker Compose, Kubernetes, registries, multi-stage builds, or deployment yet.

---

## 43. Debugging Guide

### Parser bugs

Symptoms:

- model outputs Markdown
- JSON parse fails
- missing fields

Fixes:

- stricter prompt
- show one valid example
- strip code fences
- retry once with "return valid JSON only"
- log raw output

### Tool execution bugs

Symptoms:

- command works manually but fails in runner
- file not found
- timeout

Fixes:

- print cwd
- list files before command
- cap timeouts

- use absolute tempdir internally
- never use shell-specific syntax in MVP

## **Verifier bugs**

Symptoms:

- correct solution marked wrong
- wrong solution marked right

Fixes:

- add debug output
- save workspace after failed runs
- manually rerun verifier
- add hidden checks
- hash test files

## **DPO bugs**

Symptoms:

- loss NaN
- CUDA OOM
- no improvement
- model becomes verbose/weird

Fixes:

- smaller model
- lower LR
- shorter sequences
- audit preference format
- high-margin pairs only
- SFT warmup

## **Dashboard bugs**

Symptoms:

- app slow
- JSON too large
- charts unreadable

Fixes:

- sample trajectories
  - cache loaded files
  - show only selected task
  - use pandas summaries
- 

## 44. Demo Strategy

Your demo should be founder-readable in 90 seconds.

### Demo flow

```
flowchart LR
    A["Problem"] --> B["Agent trace"]
    B --> C["Verifier score"]
    C --> D["Preference pair"]
    D --> E["DPO result"]
    E --> F["Baseline vs tuned dashboard"]
```

### Script

I built AgentAlign Lab, a verifier-guided DPO pipeline for terminal agents.

The problem is that tool-using agents often look confident while failing tests, looping, or taking unsafe actions.

My system runs small coding and data tasks, records every tool call, scores the trajectory with deterministic verifiers, and converts better and worse traces into DPO preference pairs.

I fine-tuned a small Qwen coder model with LoRA and compared it against the baseline on held-out tasks.

The dashboard shows pass rate, unsafe actions, invalid tool calls, and side-by-side chosen/rejected trajectories.

### Demo assets

Need:

- 1 GIF or short video
  - 1 screenshot of dashboard
  - 1 architecture diagram
  - 1 results table
  - 1 example preference pair
- 

## 45. README Strategy

README first screen must show:

1. What it is
2. Why it matters
3. Results
4. Demo
5. How it works

## README skeleton

```
# AgentAlign Lab

Verifier-guided DPO for reliable terminal agents.

## Result Snapshot

| Model | Pass rate | Unsafe actions | Invalid actions |
|---|---:|---:|---:|
| Base | ... | ... | ... |
| DPO-tuned | ... | ... | ... |

## Demo

Video/GIF here.

## How It Works

Diagram here.

## Quickstart

Commands here.

## Dataset

Task, trajectory, and preference schemas.

## Training

DPO/LoRA setup.

## Evaluation

Held-out task results.

## Failure Analysis

What failed and why.
```

## README rule

Do not make people scroll through theory before seeing results.

---

## 46. Resume Bullet Strategy

Use numbers. Use verbs. Use systems language.

## Strong bullets

Built AgentAlign Lab, a verifier-guided DPO pipeline for terminal agents that generated 600+ tool-use trajectories and 1,200+ preference pairs from deterministic task verifiers.

Fine-tuned Qwen2.5-Coder-1.5B with QLoRA/DPO on verifier-ranked agent traces and evaluated baseline vs tuned behavior on held-out coding/data tasks.

Designed an agent observability dashboard showing trajectory timelines, verifier scores, unsafe actions, invalid tool calls, and failure categories across model runs.

Implemented a local sandboxed terminal-agent harness with Pydantic action schemas, subprocess allowlists, pytest-based verifiers, and JSONL trace logging.

Avoid:

Worked on RLHF project.

Too vague.

---

## 47. Founder Outreach Strategy

Do not ask first. Show proof first.

### Base message

Hi <Name>, I am a 3rd-year AIML student and built AgentAlign Lab: a verifier-guided DPO pipeline that turns terminal-agent traces into preference data.

It runs coding/data tasks, logs every tool call, scores trajectories with deterministic verifiers, and fine-tunes a small open model to reduce failed or unsafe tool use.

I thought this was relevant to <Company> because you work on <specific: evals/post-training/agent tools/observability>.

Repo/demo: <link>. I would love to contribute as an AI/ML intern if this direction is useful.

### Company-specific angles

| Target                  | Emphasize                                          |
|-------------------------|----------------------------------------------------|
| Rubric/post-training    | RLAIF-style verifier feedback, DPO pairs           |
| Braintrust/Galileo/Moda | evals, traces, failure analytics                   |
| Portkey                 | observability, guardrails, LLM gateway reliability |
| Composio                | tool-use reliability and action traces             |
| Dex/Skyvern             | agent trajectories and task success                |
| E2B/Unbound/Salus       | sandboxing, safe execution, forbidden actions      |
| OpenPipe                | fine-tuning data pipeline                          |

| Target                      | Emphasize                                |
|-----------------------------|------------------------------------------|
| Sarvam                      | possible Indic eval/preference extension |
| Arintra/Yellow.ai/Richpanel | workflow-agent reliability               |

## Outreach rule

Send 20 messages:

- 5 YC agent/eval/post-training companies
- 5 Bangalore AI infra/eval companies
- 5 applied agent startups
- 5 AI/ML internship openings

Track responses in a spreadsheet.

---

## 48. Internship Positioning Strategy

Position yourself as:

AI systems and post-training engineering intern.

Not:

beginner ML student looking for any internship.

Your portfolio story:

1. I know ML/DL fundamentals.
2. I can fine-tune small models.
3. I can build agent loops and verifiers.
4. I can evaluate failures quantitatively.
5. I can ship a demo and explain tradeoffs.

Roles to apply for:

- AI/ML Intern
- AI Engineering Intern
- LLM Engineering Intern
- GenAI Intern
- Applied AI Intern
- ML Engineering Intern
- AI Infra Intern



- Evaluation Engineer Intern
- Forward Deployed AI Engineer Intern

Keywords for resume/LinkedIn:

- LLM post-training
- DPO
- QLoRA
- agent evaluation
- tool-use agents
- verifier-guided learning
- preference datasets
- AI observability
- PyTorch
- Hugging Face
- Gradio
- Pydantic
- pytest

## 49. What Makes This Research-Grade

Research-grade does not mean huge. It means disciplined.

This project becomes research-grade if it has:

- a clear research question
- train/val/test split
- baseline comparison
- ablation
- failure taxonomy
- limitations
- reproducibility
- honest negative results

### Minimum ablations

| Ablation                       | Question                        |
|--------------------------------|---------------------------------|
| all pairs vs high-margin pairs | Does preference quality matter? |

| Ablation                             | Question                                      |
|--------------------------------------|-----------------------------------------------|
| no safety penalty vs safety penalty  | Does safety scoring reduce unsafe actions?    |
| trajectory-level vs step-level pairs | Does process feedback help?                   |
| SFT warmup vs DPO only               | Does imitation help before preference tuning? |

## Research claim you can safely make

Verifier-guided preference data can improve some measurable dimensions of small terminal-agent behavior, especially structured action validity, unsafe action reduction, and task success on simple held-out tasks.

Do not claim:

This solves RLHF or alignment.

## 50. What Makes This Startup-Grade

Startup-grade means usable, measurable, and easy to inspect.

You need:

- clear README
- runnable scripts
- simple dashboard
- real metrics
- screenshots
- demo video
- concise architecture
- no broken dependencies
- no huge setup burden

Startup-grade checklist:

- ☐ README.md has result table near top.
- ☐ python scripts/07\_launch\_dashboard.py works.
- ☐ data/samples/ has small demo data.
- ☐ report/paper.md explains methods.
- ☐ configs/ includes reproducible training settings.
- ☐ failures are documented.
- ☐ limitations are honest.

- ☐ repo does not require paid API keys for the basic demo.

---

## 51. Final Recommended Execution Strategy

Build in this order:

1. Task/verifier.
2. Agent loop.
3. Traces.
4. Preference data.
5. DPO training.
6. Evaluation.
7. Dashboard.
8. Report.
9. Outreach.

Spend your effort like this:

| Area                        | Effort     |
|-----------------------------|------------|
| task/verifier/data quality  | 30 percent |
| agent runner/logging        | 20 percent |
| DPO training                | 15 percent |
| evaluation/failure analysis | 20 percent |
| dashboard/README/outreach   | 15 percent |

This is the right balance for internship success.

---

## 52. Final Reality Check

Hard truths:

- The model may not improve dramatically.
- Free GPUs will be annoying.
- Your first preference dataset will be noisy.
- Your first agent loop will break often.
- The dashboard will reveal embarrassing failures.

That is normal.

The win is not a perfect model. The win is a complete engineering loop:

```
agent behavior -> trace -> verifier -> preference data -> fine-tune -> eval -> failure analysis
```

That loop is exactly what modern AI infra, evals, and post-training teams care about.

If you execute this cleanly, you will have a stronger internship project than most students applying with generic GenAI wrappers.

---

## 53. Final "Do This First Tomorrow" Action Plan

Do these 10 things first:

1. Create repo `agentalign-lab`.
2. Add `README.md` with one-paragraph project pitch.
3. Add `pyproject.toml`.
4. Create `src/agentalign/schemas.py`.
5. Define `Task`, `Action`, `Step`, `Trajectory`, and `VerifierResult` Pydantic models.
6. Manually write 5 Python bug-fix tasks.
7. Implement `tempdir` workspace creation.
8. Implement `pytest` verifier.
9. Run verifier on all 5 tasks.
10. Commit the working skeleton.

Tomorrow's success condition:

```
I can load a task, create its workspace, run its verifier, and save a valid result JSON.
```

Do not touch DPO tomorrow.

Do not build dashboard tomorrow.

Do not read five papers tomorrow.

Build the first working verifier. Everything else grows from that.

---

## 54. V2 Upgrade Map: How To Use This Handbook Daily

This V2 upgrade does not replace the original 53 sections. Treat Sections 1-53 as the technical blueprint and Sections 54 onward as the operating system for actually building it without getting lost.

The project direction stays the same:

```
Verifier-guided preference learning for terminal agents.
```

The real upgrade is execution pressure. You are not just reading post-training theory. You are building a loop that looks like real AI systems work:

tasks -> agent attempts -> traces -> verifiers -> preference pairs -> DPO -> evals -> failure analytics -> demo

Use the handbook in this order:

| Moment                   | Section to use | Purpose                                    |
|--------------------------|----------------|--------------------------------------------|
| Starting the day         | 55             | Decide one concrete output for today       |
| Implementing a component | 56             | See what good implementation looks like    |
| Something breaks         | 57 and 68      | Debug without rewriting everything         |
| Cleaning code            | 58             | Keep the repo startup-readable             |
| Learning theory          | 59             | Learn only what helps the build            |
| Preparing outreach       | 60 and 61      | Turn the project into interview signal     |
| Choosing tools           | 62             | Know modern infra without overbuilding     |
| Explaining architecture  | 63             | Use diagrams and system flows              |
| Feeling scattered        | 64 and 65      | Reset priorities and mindset               |
| Planning the next week   | 67             | Convert ambition into weekly deliverables  |
| Final polish             | 69             | Quality-control the project like a product |

The central rule:

Every day must produce a durable artifact: code, data, eval results, a trace, a plot, a README section, or a written decision.

Learning is allowed. Wandering is not.

---

## 55. Daily Operating System

This project is too large to run on mood. It needs a daily operating system.

Your daily goal is not "work on AgentAlign." That is vague. Your daily goal is:

Ship one inspectable artifact that moves the system closer to a demo.

An inspectable artifact can be:

- a working CLI command
- a verifier that catches a real failure

- 10 new task JSON files
- 20 validated preference pairs
- an eval table
- a dashboard panel
- a trace viewer
- a bug report with root cause
- a README screenshot
- a training run summary

## The Daily Loop

```

flowchart TD
    A["Morning: choose one artifact"] --> B["Code: build the smallest working path"]
    B --> C["Run: execute on real tasks"]
    C --> D{"Passes?"}
    D -- "yes" --> E["Log result and commit"]
    D -- "no" --> F["Debug with isolation"]
    F --> G["Write failure note"]
    G --> B
    E --> H["End-of-day review"]
    H --> I["Choose tomorrow's first command"]
  
```

## Morning Workflow: 15 Minutes

Open your repo and answer these in a daily note:

```

# Daily Note: YYYY-MM-DD

## Today's one output
- I will produce:

## Why this matters
- This unlocks:

## First command I will run
~~~bash

~~~

## Risk
- The thing most likely to block me:

## Fallback
- If blocked after 60 minutes, I will:
  
```

Good daily outputs:

- "Implement `run_task --task-id python_fix_001` and save one trajectory JSON."
- "Add verifier result schema and make all verifier outputs valid Pydantic objects."
- "Generate 25 candidate/chosen/rejected pairs and manually audit 10."

- "Run base model eval on 30 tasks and produce reports/base\_eval.md."

Bad daily outputs:

- "Learn agents."
- "Read about DPO."
- "Improve code."
- "Work on dashboard."
- "Research modern AI infra."

Those are moods, not outputs.

## Coding Workflow

Use this loop for every implementation block:

1. Pick one user-facing command.
2. Implement only the code needed for that command.
3. Run it on one tiny example.
4. Save the artifact.
5. Add one test or one validation check.
6. Commit.

Example:

```
agentalign tasks validate data/tasks/python_fix_001.json
agentalign run --task-id python_fix_001 --agent baseline --out runs/dev
agentalign verify --trajectory runs/dev/python_fix_001.json
agentalign eval --split dev --agent baseline
```

If you cannot express a component as a CLI command, the component is probably too blurry.

## Debugging Workflow

When something breaks, do not jump into random edits. Use the four-box method:

Input -> Component -> Output -> Invariant

Ask:

- What exact input entered the component?
- What exact output did it produce?
- What invariant should always hold?
- Which smallest script reproduces the failure?

Debug in this order:

```
schema -> file paths -> subprocess -> timeout -> verifier -> model output ->
training
```

Most bugs in this project will not be "deep ML bugs." They will be boring systems bugs: wrong path, missing file, invalid JSON, stale tempdir, nondeterministic verifier, or silently malformed training data.

## End-Of-Day Review Workflow: 20 Minutes

End every session with this:

```
# End of Day Review: YYYY-MM-DD

## Artifact shipped
- File/command/result:

## Evidence it works
- Command:
- Output:
- Screenshot/plot/report:

## What broke
- Symptom:
- Root cause:
- Fix:

## Dataset/eval changes
- Added:
- Removed:
- Audited:

## Tomorrow's first command
~~~bash

~~~

## One thing I learned
-
```

If there is no evidence it works, it is not done.

## Commit Strategy

Commits are not just Git hygiene. They are project memory.

Commit when one of these becomes true:

- a CLI command works end to end
- a schema is stable
- a verifier is added with tests
- a dataset batch is added and validated
- an eval run produces a report
- a bug is fixed with a regression test
- the README gains a screenshot or metric

Commit message style:

```
feat(tasks): add python bug-fix task schema
feat(verifier): run pytest verifier in temp workspace
fix(parser): retry malformed action JSON once
eval(base): add baseline pass-rate report for 30 tasks
docs(readme): add architecture diagram and demo flow
```



Avoid giant commits like:

```
update project
many changes
final
work
```

Those make you look less experienced because they hide engineering decisions.

## Experiment Logging Strategy

Every experiment gets an ID:

```
exp_YYYYMMDD_short_name
```

Example:

```
exp_20260604_dpo_qwen15b_lr5e6_beta01
```

Each experiment should create:

```
experiments/
  exp_20260604_dpo_qwen15b_lr5e6_beta01/
    config.yaml
    command.txt
    train_log.txt
    metrics.json
    eval_before.json
    eval_after.json
    notes.md
```

Minimum notes.md:

```
# Experiment: exp_20260604_dpo_qwen15b_lr5e6_beta01

## Hypothesis
- DPO on verifier-derived preference pairs will improve pass rate on held-out terminal tasks.

## Setup
- Base model:
- Dataset version:
- Pair count:
- Train split:
- Eval split:
- LoRA rank:
- Learning rate:
- Beta:
- Epochs:

## Result
- Before pass rate:
- After pass rate:
- Invalid action rate:
- Avg steps:
- Timeout rate:

## Interpretation
-

## Decision
- Keep / rerun / discard / investigate
```

## Research-To-Implementation Workflow

Never read a paper without extracting implementation obligations.

Use this template:

```
# Reading Note

## Resource
- Link:

## Why I am reading it
-

## 3 ideas relevant to AgentAlign
1.
2.
3.

## 1 thing I will implement
-

## 1 thing I will ignore for MVP
-

## Code/file affected
-
```

Example:

```
DPO paper -> implement chosen/rejected pairs and beta sweep.
ReAct paper -> implement step-level action/observation traces.
Terminal-Bench -> design tasks as isolated terminal environments with verifiers.
Phoenix/Langfuse docs -> log trace spans, but defer full hosted observability.
```

## Daily Notes Template

```
# AgentAlign Daily Note: YYYY-MM-DD

## Target artifact
-

## Today's scope
- Must:
- Should:
- Could:

## Current system state
- Tasks:
- Verifiers:
- Agent:
- Preference pairs:
- Eval:
- Training:
- Dashboard:

## Commands run
~~~bash

~~~

## Result
-
```

```
## Bugs found
-

## Decisions made
-

## Tomorrow
-
```

## Debugging Log Template

```
# Debug Log: bug_YYYYMMDD_short_name

## Symptom
-

## Reproduction command
~~~bash

~~~

## Expected
-

## Actual
-

## Suspected causes
1.
2.
3.

## Isolation steps
- [ ] Validate input schema
- [ ] Print absolute paths
- [ ] Run component directly
- [ ] Reduce to one task
- [ ] Check logs
- [ ] Add regression test

## Root cause
-

## Fix
-

## Prevention
-
```

## Architecture Decision Template

```
# ADR: short-title

## Date
YYYY-MM-DD

## Status
Accepted / Rejected / Revisit later

## Context
-
```

```
## Decision
-

## Alternatives considered
-

## Consequences
-

## Startup relevance
-
```

Example ADRs:

- use JSONL for trajectories instead of SQLite in MVP
- use deterministic verifiers instead of LLM judges in MVP
- use Gradio dashboard instead of React
- defer Docker sandbox to V2
- use Qwen2.5-Coder-1.5B for Colab DPO

## Failed Experiment Template

Failed experiments are valuable if they are legible.

```
# Failed Experiment: exp_id

## What I tried
-

## Why I expected it to work
-

## What happened
-

## Evidence
-

## Root cause category
- Data / model / verifier / training / eval / infra / unknown

## What I learned
-

## What I will do differently
-
```

## Learning Summary Template

```
# Weekly Learning Summary

## Concepts learned
-

## Systems lessons
-
```

```
## ML lessons
-

## Debugging lessons
-

## Interview talking points created
-

## Next week's risk
-
```

## Two-Hour Productive Session Template

If you only have two hours, do this:

| Time        | Action                              | Output               |
|-------------|-------------------------------------|----------------------|
| 0-10 min    | Pick one artifact                   | Written target       |
| 10-25 min   | Run current failing or next command | Baseline evidence    |
| 25-75 min   | Implement smallest change           | Code                 |
| 75-95 min   | Run on one real task                | Result file          |
| 95-110 min  | Add validation/test/log             | Safety check         |
| 110-120 min | Commit and write next command       | Clean stopping point |

Two-hour example:

```
Artifact: verifier writes structured failure reasons.
Command: agentalign verify --task python_fix_001
Done when: verifier_result.json contains passed, score, failure_tags, stdout,
stderr, duration_ms.
```

## Minimum Daily Progress Definition

On a bad day, minimum progress is:

Run one command, inspect one failure, write one note, make one tiny improvement.

Examples:

- Fix one schema validation error.
- Add one missing timeout.
- Write one task.
- Audit five preference pairs.
- Add one README metric.
- Reproduce one bug and write the log.

This prevents the project from becoming all-or-nothing.

## What To Do When Overwhelmed

When overwhelmed, your brain will try to solve the whole project at once. Refuse.

Do this:

1. Open README .md.
2. Write the current one-sentence project status.
3. Run the last known working command.
4. Pick the nearest broken component.
5. Reduce to one task and one model output.
6. Write the exact failing command.
7. Fix only that.

Use this reset statement:

```
I am not solving alignment today. I am making one artifact more reliable.
```

The hierarchy of recovery:

```
working verifier > working agent > more tasks > preference pairs > DPO >
dashboard > polish
```

If in doubt, go back to the verifier.

## How To Avoid Tutorial Hell

Tutorial hell happens when reading feels safer than building.

Your rule:

```
For every 45 minutes of reading, produce 45 minutes of code or notes tied to the
repo.
```

Allowed reading:

- read TRL DPO docs while implementing `train_dpo.py`
- read Docker run docs while isolating task execution
- read Phoenix docs while designing trace schema
- read ReAct while implementing action/observation loops

Avoidance reading:

- comparing every agent framework
- watching "RLHF full course" before writing one preference pair
- reading every DPO variant before training a LoRA
- studying Kubernetes for a solo Colab project
- rewriting the stack because a blog said a new framework is better

Use this filter:

Will this resource change today's code, dataset, eval, or demo?

If no, save it for V2.

---

## 56. Implementation Deep Dives

This section is the senior-engineer companion for each major component. Use it before implementing, during debugging, and before calling a component "done."

The system has twelve core components:

```
flowchart LR
    A["Task generation"] --> B["Verifier design"]
    B --> C["Agent loop"]
    C --> D["Tool execution"]
    D --> E["JSON parsing and retries"]
    E --> F["Trajectory logging"]
    F --> G["Preference pair generation"]
    G --> H["DPO dataset quality"]
    H --> I["DPO training"]
    I --> J["Evaluation"]
    J --> K["Failure analysis"]
    K --> L["Dashboard"]
    L --> A
```

### 56.1 Task Generation Deep Dive

Implementation philosophy:

Tasks are not prompts. Tasks are testable environments with ground-truth verification.

A good task has:

- a clear goal
- a small workspace
- deterministic verifier
- no hidden internet dependency
- meaningful failure modes
- expected difficulty label
- tags for later analysis

Bad task:

"Fix the code."

Good task:

```
{
  "task_id": "py_fix_001",
  "category": "python_bugfix",
  "prompt": "Fix the failing tests in this small Python package.",
}
```

```

"workspace_template": "templates/py_fix_001",
"verifier": {
  "type": "pytest",
  "command": "pytest -q",
  "timeout_sec": 20
},
"difficulty": "easy",
"tags": ["pytest", "off_by_one", "single_file"]
}

```

Common mistakes:

- making tasks too broad
- mixing multiple bugs in one task early
- relying on natural language success instead of tests
- creating tasks that pass without changes
- creating tasks that are impossible because setup is broken
- forgetting to version tasks

Debugging checklist:

- Does the workspace copy correctly?
- Does the verifier fail before the fix?
- Does the verifier pass after the reference fix?
- Can the task run offline?
- Is the timeout enough but not infinite?
- Are stdout and stderr captured?
- Does the task have one intended learning signal?

Engineering tradeoffs:

| Choice               | Pros                       | Cons                   | MVP decision |
|----------------------|----------------------------|------------------------|--------------|
| Handwritten tasks    | High quality, controllable | Slower to scale        | Use for MVP  |
| Synthetic LLM tasks  | Fast volume                | Noisy, unverifiable    | V2 only      |
| Real GitHub issues   | Impressive                 | Setup pain, huge scope | Defer        |
| Tiny unit-test tasks | Deterministic              | Less flashy            | Best early   |

What good code looks like:

- task schema validated by Pydantic
- task files live under versioned directories
- task loading has explicit error messages
- verifier command stored in task metadata



- task generation separated from task execution

What bad code looks like:

- task prompt embedded inside runner code
- hardcoded paths like `../..../tasks`
- verifier logic hidden in random scripts
- no reference solution
- no task validation command

Startup parallel:

This is exactly like eval-set design at an AI startup. The task suite is the product's truth source. If the evals are sloppy, every model improvement claim becomes suspicious.

Why decisions matter:

DPO quality depends on task quality. Bad tasks create bad preference pairs. Bad preference pairs create bad fine-tuning. Bad fine-tuning creates fake progress.

## 56.2 Verifier Design Deep Dive

Implementation philosophy:

The verifier is the judge. It should be deterministic, boring, and hard to fool.

Verifier output should always be structured:

```
{
  "task_id": "py_fix_001",
  "passed": false,
  "score": 0.0,
  "failure_tags": ["pytest_failed"],
  "stdout": "...",
  "stderr": "...",
  "duration_ms": 1834
}
```

Verifier pipeline:

```
flowchart TD
    A["Trajectory final workspace"] --> B["Run verifier command"]
    B --> C["Capture stdout/stderr/exit code"]
    C --> D["Parse result"]
    D --> E["Assign score and failure tags"]
    E --> F["Validate VerifierResult schema"]
    F --> G["Persist verifier_result.json"]
```

Common mistakes:

- only returning pass/fail
- not recording failure reason
- no timeout
- no stderr capture

- relying on fragile string matching
- letting verifier mutate the original task template
- not testing the verifier itself

Debugging checklist:

- Does verifier fail on initial broken workspace?
- Does verifier pass on reference solution?
- Does timeout return a structured failure?
- Are relative paths resolved from the task workspace?
- Does verifier clean up after itself?
- Does score mean the same thing across tasks?

Engineering tradeoffs:

| Verifier type | Use when             | Risk                        |
|---------------|----------------------|-----------------------------|
| Pytest        | Python bug-fix tasks | tests may be too narrow     |
| Shell command | file/output tasks    | brittle command assumptions |
| Static checks | format/config tasks  | can reward shallow fixes    |
| LLM judge     | semantic tasks       | noisy and expensive         |
| Hybrid        | V2 credibility       | more complexity             |

What good code looks like:

```
def run_verifier(task: Task, workspace: Path) -> VerifierResult:
    started = time.monotonic()
    try:
        proc = subprocess.run(
            task.verifier.command,
            cwd=workspace,
            shell=False,
            capture_output=True,
            text=True,
            timeout=task.verifier.timeout_sec,
        )
        return VerifierResult.from_process(task.task_id, proc, started)
    except subprocess.TimeoutExpired as exc:
        return VerifierResult.timeout(task.task_id, exc, started)
```

What bad code looks like:

```
os.system("pytest")
return True
```

Startup parallel:

Verifier design maps directly to evals engineering. Startups care because deterministic evals let teams decide whether a model, prompt, or product change actually improved reliability.

Why decisions matter:

Your verifier decides which trajectories become "chosen" and "rejected." If the verifier lies, DPO learns the lie.

## 56.3 Agent Loop Architecture Deep Dive

Implementation philosophy:

Keep the agent loop explicit. Do not hide everything inside an agent framework for MVP.

The loop:

```
sequenceDiagram
    participant Runner
    participant Model
    participant Parser
    participant Tool
    participant Log

    Runner->>Model: prompt + task + history
    Model-->>Runner: proposed action JSON
    Runner->>Parser: validate action
    alt valid action
        Parser-->>Runner: Action object
        Runner->>Tool: execute action
        Tool-->>Runner: observation
    else invalid action
        Parser-->>Runner: parser error
        Runner->>Model: retry with error feedback
    end
    Runner->>Log: append step
    Runner->>Runner: stop or continue
```

Common mistakes:

- unlimited loops
- no max step count
- no stop action
- not logging failed parse attempts
- feeding huge observations back into the prompt
- not separating model call from tool execution
- no reproducible seed/config

Debugging checklist:

- Can the loop run with a fake model?
- Can it stop after `max_steps`?
- Does every step have action, observation, timestamps, and errors?
- Are invalid actions counted?
- Is the final workspace saved or inspectable?
- Can you replay the trajectory?

Engineering tradeoffs:

| Design               | Pros                      | Cons                    | MVP decision    |
|----------------------|---------------------------|-------------------------|-----------------|
| Custom loop          | Full control, educational | More code               | Use             |
| LangGraph            | Robust workflows          | More abstraction        | Study, V2       |
| OpenAI Agents SDK    | Excellent reference       | API/provider dependency | Study, not core |
| AutoGen/CrewAI style | Quick demos               | Less eval-control       | Avoid for MVP   |

What good code looks like:

- `AgentRunner` owns orchestration
- `ModelClient` owns generation
- `ActionParser` owns validation
- `ToolExecutor` owns subprocess/file operations
- `TrajectoryLogger` owns persistence

What bad code looks like:

- one 500-line function
- prompt construction mixed with subprocess execution
- exceptions swallowed
- no typed action schema
- no replay mode

Startup parallel:

Agent/tool startups care about controllability. A simple transparent loop beats a magical framework when you are debugging reliability.

Why decisions matter:

Your final demo is not "the model answered." It is "we can explain why the agent failed and improve it through preference learning."

## 56.4 Tool Execution Layer Deep Dive

Implementation philosophy:

Tool execution is untrusted I/O. Treat every command as a risky operation even in a toy project.

Low-level flow:

```
flowchart TD
    A["Action object"] --> B{"Allowed tool?"}
    B -- "no" --> C["Reject action"]
    B -- "yes" --> D["Resolve workspace path"]
```

```
D --> E{"Within workspace?"}  
E -- "no" --> C  
E -- "yes" --> F["Run subprocess with timeout"]  
F --> G["Capture stdout/stderr/exit"]  
G --> H["Truncate observation"]  
H --> I["Return ToolResult"]
```

Common mistakes:

- `shell=True` by default
- no command allowlist
- no timeout
- no output truncation
- allowing writes outside workspace
- not capturing exit code
- treating failed command as Python exception only

Debugging checklist:

- Print `cwd` before execution.
- Print resolved absolute paths.
- Log exit code.
- Log timeout status.
- Limit output size.
- Test command success and failure.
- Test command timeout.

Engineering tradeoffs:

| Choice                | Benefit          | Risk             |
|-----------------------|------------------|------------------|
| Local subprocess      | Simple and fast  | weaker isolation |
| Docker container      | better isolation | setup cost       |
| Restricted shell tool | transparent      | less realistic   |
| Full terminal agent   | impressive       | dangerous scope  |

MVP decision:

Use local `tempdir` execution with strict path checks and timeouts. Add Docker in V2 only after the rest works.

Startup parallel:

This maps to tool safety, sandboxing, agent guardrails, and production blast-radius control. Founders will care that you considered execution risk.

## 56.5 JSON Parsing And Retries Deep Dive

Implementation philosophy:

Invalid JSON is not an edge case. It is a core agent failure mode.

Parser lifecycle:

```
flowchart TD
    A["Raw model output"] --> B["Extract candidate JSON"]
    B --> C{"Valid JSON?"}
    C -- "no" --> D["Parser error object"]
    C -- "yes" --> E["Pydantic validation"]
    E --> F{"Valid schema?"}
    F -- "no" --> D
    F -- "yes" --> G["Action object"]
    D --> H{"Retries left?"}
    H -- "yes" --> I["Send concise repair prompt"]
    H -- "no" --> J["Record invalid_action failure"]
```

Common mistakes:

- regex-only parsing with no schema validation
- unlimited retries
- not saving raw invalid output
- overcomplicated JSON repair
- accepting unknown action fields silently

Debugging checklist:

- Save raw model output.
- Save parser error message.
- Validate with Pydantic.
- Count invalid action rate.
- Test with malformed examples.
- Add one repair retry, not five.

Good parser behavior:

- strict schema
- one repair retry
- clear error returned to model
- invalid outputs logged as steps
- invalid action rate included in eval metrics

Bad parser behavior:

- silently inventing missing fields
- converting any text into best-effort commands

- ignoring extra fields
- dropping parse failures from trajectories

Startup parallel:

Structured outputs, constrained decoding, and schema validation are now basic production AI engineering. A startup wants to see that you can make LLM outputs machine-usable.

## 56.6 Trajectory Logging Deep Dive

Implementation philosophy:

If it is not logged, it did not happen. Trajectories are the core dataset for evals, preference learning, and failure analysis.

Minimum trajectory schema:

```
{
  "trajectory_id": "run_001_py_fix_001_baseline",
  "task_id": "py_fix_001",
  "agent_id": "baseline_qwen",
  "started_at": "2026-06-01T10:00:00Z",
  "ended_at": "2026-06-01T10:00:22Z",
  "steps": [],
  "verifier_result": {},
  "metadata": {
    "model": "Qwen/Qwen2.5-Coder-1.5B-Instruct",
    "temperature": 0.2,
    "max_steps": 8,
    "code_version": "git_sha"
  }
}
```

Trace lifecycle:

```
flowchart TD
    A["Run config"] --> B["Trajectory initialized"]
    B --> C["Step appended after each action"]
    C --> D["Verifier result attached"]
    D --> E["Trajectory schema validated"]
    E --> F["JSONL persisted"]
    F --> G["Eval reads trajectory"]
    F --> H["Dashboard reads trajectory"]
    F --> I["Preference generator reads trajectory"]
```

Common mistakes:

- saving only final answer
- no timestamps
- no model config
- not logging failed actions
- no raw model output
- no git/version metadata
- overwriting runs

Debugging checklist:

- Is every run ID unique?
- Can you load every saved trajectory?
- Do failed runs save too?
- Is raw model output stored?
- Are observations truncated but recoverable?
- Can dashboard render the trace?

Startup parallel:

This is observability. LLM teams need traces because aggregate metrics do not explain failures. A strong trace schema makes your project feel real.

## 56.7 Preference Pair Generation Deep Dive

Implementation philosophy:

Preference pairs are not just winners and losers. They are behavioral examples teaching the model what reliable tool use looks like.

Pair generation flow:

```
flowchart TD
    A["Multiple trajectories for same task"] --> B["Group by task_id"]
    B --> C["Rank by verifier score"]
    C --> D["Select chosen trajectory"]
    C --> E["Select rejected trajectory"]
    D --> F["Build DPO prompt"]
    E --> F
    F --> G["Validate pair schema"]
    G --> H["Manual audit sample"]
    H --> I["Save preference JSONL"]
```

Common mistakes:

- pairing trajectories from different tasks
- pairing two equally bad runs
- using verifier pass/fail only when partial scores exist
- no manual audit
- pairs where chosen and rejected differ only by formatting
- too many near-duplicate pairs

Debugging checklist:

- Same task ID?
- Same prompt?
- Chosen score greater than rejected score?
- Meaningful behavioral difference?



- No leaked verifier output in prompt?
- No broken JSON in chosen/rejected?
- Audit at least 10-20 percent early.

Good pair:

Prompt: terminal task instructions + available tools.

Chosen: concise actions that inspect files, modify bug, run tests, stop.

Rejected: loops, ignores test failure, edits wrong file, invalid JSON, or times out.

Bad pair:

Chosen and rejected both fail with same error.

Startup parallel:

Preference data quality is one of the highest-leverage post-training skills. Startups know models often improve less from algorithm novelty and more from clean data.

## 56.8 DPO Dataset Quality Deep Dive

Implementation philosophy:

The DPO dataset is a product artifact. Treat it like code.

Dataset quality dimensions:

| Dimension       | Good                                     | Bad                        |
|-----------------|------------------------------------------|----------------------------|
| Correctness     | chosen behavior is actually better       | chosen only looks better   |
| Diversity       | multiple task types/failures             | same task repeated         |
| Contrast        | clear difference between chosen/rejected | weak preference signal     |
| Formatting      | consistent prompt/completion format      | mixed templates            |
| Safety          | no accidental path leaks/secrets         | raw local paths everywhere |
| Split integrity | train/test separate tasks                | same task appears in both  |

Common mistakes:

- training on eval tasks
- not freezing dataset version
- mixing schemas
- no pair statistics
- no invalid pair filter
- too few high-quality pairs

Debugging checklist:

- Count pairs per task.
- Count pairs per failure tag.
- Check chosen score distribution.
- Check rejected score distribution.
- Ensure no eval task IDs in train.
- Validate token lengths.
- Print 20 random pairs.

Startup parallel:

This mirrors data curation at alignment/post-training startups. The ability to discuss dataset quality is a major interview signal.

## 56.9 DPO Training Deep Dive

Implementation philosophy:

DPO is not magic. It amplifies the preferences you feed it. Keep the first run tiny and measurable.

Training flow:

```
flowchart TD
    A["Base model"] --> B["Load LoRA config"]
    C["Preference JSONL"] --> D["Tokenize chosen/rejected"]
    B --> E["TRL DPOTrainer"]
    D --> E
    E --> F["Train adapter"]
    F --> G["Save adapter + config"]
    G --> H["Run same eval before/after"]
    H --> I["Write experiment report"]
```

Common mistakes:

- too high learning rate
- too many epochs on tiny data
- no held-out eval
- no base-model comparison
- no sample outputs before/after
- training on noisy pairs
- assuming lower loss means better agent behavior

Debugging checklist:

- Can you overfit 10 pairs intentionally?
- Are chosen/rejected labels in correct order?
- Are prompts identical for each pair?

- Is adapter actually loaded during eval?
- Did you freeze base model?
- Is eval deterministic enough?
- Is improvement larger than noise?

Recommended first DPO run:

| Setting    | Value                                          |
|------------|------------------------------------------------|
| Model      | Qwen2.5-Coder-1.5B-Instruct                    |
| Method     | LoRA/QLoRA                                     |
| Pair count | 100-300                                        |
| Epochs     | 1                                              |
| LR         | 5e-6 to 2e-5                                   |
| Beta       | 0.05 to 0.1                                    |
| Eval tasks | held-out 20-50                                 |
| Claim      | modest pass-rate or invalid-action improvement |

Startup parallel:

Post-training teams care that you can connect training to evals. "I trained DPO" is weak. "I trained DPO, audited data, compared pass rate, and found it reduced invalid tool calls by X" is strong.

## 56.10 Evaluation Methodology Deep Dive

Implementation philosophy:

Evaluation is the truth layer. It must compare systems under the same conditions.

Evaluation flow:

```

flowchart TD
    A["Eval split"] --> B["Run base agent"]
    A --> C["Run tuned agent"]
    B --> D["Collect trajectories"]
    C --> E["Collect trajectories"]
    D --> F["Verifier metrics"]
    E --> F
    F --> G["Aggregate by task type"]
    G --> H["Failure tag comparison"]
    H --> I["Report with confidence caveats"]
  
```

Common mistakes:

- changing prompt and model at same time
- evaluating on training tasks
- reporting only cherry-picked examples
- no failure breakdown

- no latency/cost/step metrics
- no invalid action metric

Debugging checklist:

- Same task split?
- Same max steps?
- Same model temperature?
- Same verifier?
- Same timeout?
- Same prompt template except trained adapter?
- Random seed recorded?

Startup parallel:

This is what evals companies and AI infra teams obsess over. A credible evaluation methodology makes even a small project feel serious.

## 56.11 Failure Analysis Deep Dive

Implementation philosophy:

Failures are not embarrassing. Unclassified failures are embarrassing.

Failure taxonomy:

```
invalid_action
tool_error
timeout
wrong_file
insufficient_inspection
test_not_run
partial_fix
over_editing
looping
hallucinated_command
verifier_error
environment_error
```

Failure analysis workflow:

```
flowchart TD
    A["Failed trajectories"] --> B["Assign failure tags"]
    B --> C["Group by task type"]
    C --> D["Find top 3 failure modes"]
    D --> E["Pick one intervention"]
    E --> F["Prompt/data/verifier/training fix"]
    F --> G["Rerun eval"]
```

Common mistakes:

- reading failures randomly
- not tagging failures

- fixing one anecdote instead of a common pattern
- blaming model before checking task/verifier
- not distinguishing agent failure from verifier failure

Startup parallel:

This is operational AI engineering. Teams win by turning vague model badness into specific failure modes and targeted interventions.

## 56.12 Dashboard Design Deep Dive

Implementation philosophy:

The dashboard is not decoration. It is a failure analysis console.

Dashboard data flow:

```

flowchart LR
    A["tasks.jsonl"] --> D["Dashboard loader"]
    B["trajectories.jsonl"] --> D
    C["eval_results.json"] --> D
    D --> E["Metric cards"]
    D --> F["Failure table"]
    D --> G["Trace viewer"]
    D --> H["Before/after comparison"]
  
```

Good dashboard panels:

- pass rate before/after
- invalid action rate
- timeout rate
- average steps
- failure tag distribution
- task table
- trace viewer
- selected trajectory replay
- chosen/rejected pair inspector

Bad dashboard panels:

- giant hero section
- vague "AI magic" text
- charts without denominators
- only successful examples
- no failure drilldown

Startup parallel:

Evals/observability companies sell clarity. Your dashboard should show you understand why visibility matters.

## 57. Massive Debugging Playbook

Debugging this project well is part of the project. Your strongest interview stories may come from failures you diagnosed cleanly.

Use this default incident flow:

```
flowchart TD
    A["Bug appears"] --> B["Freeze reproduction command"]
    B --> C["Reduce to one task"]
    C --> D["Classify failure layer"]
    D --> E["Inspect raw artifact"]
    E --> F["Form one hypothesis"]
    F --> G["Run one test"]
    G --> H["Confirmed?"]
    H -- "yes" --> I["Patch root cause"]
    H -- "no" --> F
    I --> J["Add regression check"]
    J --> K["Write debug note"]
```

Failure layers:

| Layer          | Examples                                                |
|----------------|---------------------------------------------------------|
| Schema         | invalid task JSON, invalid action, missing field        |
| Filesystem     | wrong path, workspace not copied, stale tempdir         |
| Tool execution | subprocess error, timeout, bad cwd                      |
| Agent loop     | infinite loop, no stop action, repeated commands        |
| Verifier       | false positive, false negative, flaky tests             |
| Data           | noisy pairs, train/eval leakage, duplicate examples     |
| Training       | CUDA OOM, unstable DPO loss, adapter not loaded         |
| Eval           | inconsistent config, wrong split, cherry-picked metrics |
| Dashboard      | broken load, stale cache, wrong aggregation             |

### Invalid JSON Debugging

Symptoms:

- model output cannot parse
- parser throws `JSONDecodeError`
- action field missing
- model surrounds JSON with prose

- invalid action rate is high

Probable causes:

- prompt does not specify exact schema
- examples are inconsistent
- model is too small or temperature too high
- action schema is too complex
- parser expects one format while prompt shows another

Exact fixes:

- reduce temperature
- show one canonical JSON example
- add `Action` Pydantic schema
- allow one repair retry
- log raw output
- count invalid action rate as a metric

Prevention:

- keep action schema small
- validate every action before execution
- write parser unit tests with malformed outputs
- do not silently repair unknown fields

## Parser Failure Debugging

Symptoms:

- model output contains valid-looking JSON but schema rejects it
- actions fail because of wrong field names
- parser works for one model but not another

Probable causes:

- inconsistent key names
- Pydantic model too permissive or too strict
- model emits list instead of object
- action type enum unclear

Exact fixes:

- print parser error to debug log

- include allowed actions in prompt
- add examples for each action type
- use `extra="forbid"` so unknown fields are visible
- create `tests/test_parser.py`

Prevention:

- version your action schema
- add a CLI command: `agentalign parse --input sample.txt`
- store raw and parsed action in trajectory

## Subprocess Issues

Symptoms:

- command works manually but not in agent
- verifier cannot find files
- `FileNotFoundError`
- empty stdout/stderr
- return code unexpected

Probable causes:

- wrong cwd
- command passed as string while `shell=False`
- missing executable
- relative path resolved from repo instead of workspace
- environment variable missing

Exact fixes:

- log absolute cwd
- log exact command list
- use `shutil.which` for executables if needed
- avoid `shell=True` unless explicitly justified
- run verifier command manually inside workspace

Prevention:

- centralize subprocess wrapper
- always capture stdout, stderr, exit code, duration
- use timeouts on every subprocess call



## Timeout Debugging

Symptoms:

- agent hangs
- eval never completes
- subprocess never returns
- dashboard shows missing runs

Probable causes:

- command waiting for input
- tests stuck in infinite loop
- agent repeats same action
- no timeout at model/tool/verifier layer
- long model generation

Exact fixes:

- add timeouts at all layers
- kill process group on timeout
- record timeout failure tag
- reduce max steps
- truncate observations

Prevention:

- every external call needs a timeout
- CI/dev eval should use tiny split
- record `duration_ms` for every step

## Verifier False Positives

Symptoms:

- verifier passes a bad solution
- agent deletes tests and passes
- superficial output satisfies check

Probable causes:

- tests too weak
- verifier checks exit code only
- task allows modifying test files

- expected output check too broad

Exact fixes:

- protect test files or restore tests before verifier
- add hidden tests for key behavior
- inspect diff before verifier
- add static check for forbidden edits

Prevention:

- run verifier on known bad solutions
- include reference failing cases
- tag tasks vulnerable to shallow hacks

## **Verifier False Negatives**

Symptoms:

- correct-looking solution fails
- verifier brittle to formatting
- task passes manually but fails in runner

Probable causes:

- environment mismatch
- path assumptions
- timeout too short
- overly strict output comparison
- missing dependency

Exact fixes:

- compare manual and runner environments
- log cwd, command, stdout, stderr
- increase timeout only if justified
- make output checks semantic where possible

Prevention:

- verifier unit tests
- reference solution must pass
- avoid hidden dependency on local machine state

## **DPO Instability**

Symptoms:

- loss spikes or becomes NaN
- tuned model gets worse
- outputs become repetitive
- model stops following JSON schema
- eval metrics degrade despite lower training loss

Probable causes:

- learning rate too high
- beta poorly chosen
- noisy pairs
- too many epochs
- tokenization mismatch
- chosen/rejected swapped
- prompts too long

Exact fixes:

- lower LR
- train 1 epoch
- audit pairs
- overfit 10 clean pairs as sanity check
- verify chosen/rejected order
- shorten prompts
- evaluate invalid action rate separately

Prevention:

- start with tiny run
- save before/after samples
- freeze dataset version
- track config exactly

## **CUDA OOM**

Symptoms:

- CUDA out of memory
- notebook crashes

- training starts then dies
- eval works but training fails

Probable causes:

- model too large
- sequence length too high
- batch size too high
- no gradient checkpointing
- full fine-tuning instead of LoRA/QLoRA
- memory fragmentation

Exact fixes:

- use 1.5B-3B model
- reduce max sequence length
- batch size 1 with gradient accumulation
- use 4-bit quantization where available
- enable gradient checkpointing
- clear GPU cache between runs
- restart runtime after failed large run

Prevention:

- record GPU type and memory
- estimate tokens before training
- run a 5-step smoke test before full run

## **Bad Preference Pairs**

Symptoms:

- DPO does not improve behavior
- tuned model learns weird shortcuts
- chosen and rejected look similar
- many pairs are both failures

Probable causes:

- verifier score too coarse
- no manual audit
- weak contrast
- task duplicates

- rejected output is invalid in a way unrelated to task

Exact fixes:

- filter pairs by score margin
- require different failure tags
- manually inspect random sample
- remove ambiguous pairs
- create more high-quality handwritten tasks

Prevention:

- pair quality report
- chosen/rejected score histograms
- hold out entire tasks, not just pairs

## Overfitting Signs

Symptoms:

- train behavior improves but eval worsens
- model repeats training task phrases
- pass rate improves only on seen tasks
- invalid JSON increases on new tasks

Probable causes:

- dataset too small
- too many epochs
- train/eval leakage
- narrow task diversity

Exact fixes:

- reduce epochs
- add held-out tasks
- add more diverse pairs
- report train and eval separately

Prevention:

- split by task ID before pair generation
- keep eval set frozen
- do not tune prompts after seeing eval failures unless logged

## Noisy Data Symptoms

Symptoms:

- preference labels contradict verifier
- chosen sometimes worse than rejected
- model learns inconsistent action style

Probable causes:

- verifier bugs
- pair generator bug
- mixed prompt templates
- duplicate corrupted trajectories

Exact fixes:

- validate all pairs with schema
- recompute verifier scores from trajectory IDs
- audit pair examples by failure tag
- remove pairs with small score margin

Prevention:

- dataset build script should produce statistics and warnings
- never manually edit JSONL without validation

## Poor Model Behavior After Tuning

Symptoms:

- tuned model is more verbose
- stops using tools
- calls wrong tool
- ignores JSON schema
- fewer tasks pass

Probable causes:

- DPO pairs reward final answers, not tool correctness
- prompt format mismatch between train and eval
- overfitting
- low-quality chosen outputs
- adapter not merged/loaded correctly

Exact fixes:

- ensure train and eval prompt templates match
- include action-format examples in chosen outputs
- evaluate parser metrics, not just pass rate
- load adapter in a minimal sample script and compare outputs

Prevention:

- before/after qualitative samples
- eval on frozen split
- keep adapter config with experiment

## Hallucinated Tool Calls

Symptoms:

- model calls tools that do not exist
- command refers to nonexistent files
- model invents APIs

Probable causes:

- tool list unclear
- prompt too long
- task asks too much
- no schema enum
- no observation feedback

Exact fixes:

- use enum for tool names
- include available tools in every prompt
- reject unknown tool calls with clear error
- add failure tag `hallucinated_tool`

Prevention:

- keep tool set small
- teach from rejected examples with hallucinated tools
- include tool availability in DPO prompt

## Looping Agents

Symptoms:

- repeats same command

- edits same file repeatedly
- runs tests forever
- never emits stop action

Probable causes:

- no max step limit
- no repeated-action detection
- observation not informative
- stop criteria unclear

Exact fixes:

- enforce `max_steps`
- detect repeated action N times
- add `stop` action schema
- include last failure summary in prompt
- tag loops in failure analysis

Prevention:

- trajectory viewer should surface repeated actions
- DPO rejected examples should include loops

## **Broken Trajectory Logs**

Symptoms:

- dashboard cannot load runs
- JSONL parse error
- missing verifier results
- preference generator skips runs

Probable causes:

- partial writes
- schema changed without migration
- failed runs not finalized
- multiple processes writing same file

Exact fixes:

- write to temp file then rename
- validate trajectory before append



- include schema version
- store one JSON per run during dev

Prevention:

- add agentalign trajectories validate
- never let logging failure crash silently
- keep raw artifacts under runs/

## Top 20 Bugs You Will Almost Certainly Hit

| #  | Bug                         | Fast diagnosis         | Fix                       |
|----|-----------------------------|------------------------|---------------------------|
| 1  | Wrong workspace path        | print absolute cwd     | pass workspace explicitly |
| 2  | Verifier passes broken task | run initial template   | strengthen tests          |
| 3  | Reference solution fails    | run verifier manually  | fix task before agent     |
| 4  | Invalid JSON                | inspect raw output     | schema example + retry    |
| 5  | Model loops                 | repeated actions       | max steps + loop tag      |
| 6  | Command hangs               | no timeout             | timeout wrapper           |
| 7  | Empty stdout                | wrong command list     | log command and cwd       |
| 8  | Pair labels swapped         | chosen score lower     | assert score margin       |
| 9  | Train/eval leakage          | same task IDs          | split by task             |
| 10 | DPO worsens model           | noisy pairs or high LR | audit + lower LR          |
| 11 | OOM in Colab                | long context/model     | reduce model/seq/batch    |
| 12 | Adapter not loaded          | same outputs as base   | adapter smoke test        |
| 13 | Dashboard stale             | cached file            | reload from run path      |
| 14 | Metrics inflated            | cherry-picked split    | frozen eval split         |
| 15 | False negative verifier     | brittle check          | compare manual vs runner  |
| 16 | Tool hallucination          | unknown action         | enum + reject             |
| 17 | Broken JSONL                | partial writes         | temp write then rename    |
| 18 | Huge observations           | prompt overflow        | truncate output           |
| 19 | Flaky task                  | nondeterministic test  | remove randomness         |
| 20 | Scope creep                 | too many features      | cut to demo-critical path |

---

## 58. Engineering Best Practices

The goal is not enterprise bureaucracy. The goal is to make the repo look like a serious AI systems project built by someone who can join a startup team.

### Production-Grade Python Practices

Use:

- type hints for public functions
- Pydantic for external data schemas
- `pathlib.Path` for paths
- `subprocess.run` through one wrapper
- structured logging
- `pyproject.toml`
- `ruff` for lint/format
- `pytest` for tests
- clear CLI commands

Avoid:

- notebooks as the main codebase
- global mutable config
- hardcoded absolute paths
- `silent except Exception`
- stringly typed dictionaries everywhere
- one giant script called `main.py`

Good project boundary:

```
notebook = exploration
src/ = product code
scripts/ = reproducible commands
reports/ = results
```

### Logging Standards

Every run should log:

- run ID
- task ID
- agent ID
- model name
- temperature

- max steps
- start/end time
- tool calls
- parser failures
- verifier result
- failure tags

Log levels:

| Level   | Use                                   |
|---------|---------------------------------------|
| DEBUG   | raw model outputs, parser details     |
| INFO    | run start/end, task result            |
| WARNING | invalid action, timeout, skipped pair |
| ERROR   | component failure, corrupted artifact |

Example:

```
logger.info(
    "run_finished",
    extra={
        "run_id": run_id,
        "task_id": task.task_id,
        "passed": verifier.passed,
        "score": verifier.score,
        "steps": len(trajecory.steps),
    },
)
```

## Config Management

Keep configs explicit:

```
model:
  name: Qwen/Qwen2.5-Coder-1.5B-Instruct
  temperature: 0.2
  max_new_tokens: 512

agent:
  max_steps: 8
  parser_retries: 1

execution:
  timeout_sec: 20
  max_observation_chars: 4000

eval:
  split: dev
  task_limit: 30
```

Rules:

- config files live in `configs/`

- experiment copies config into its experiment directory
- CLI override is allowed but must be logged
- never rely on notebook state for experiments

## Folder Organization

Startup-readable structure:

```
agentalign-lab/
  README.md
  pyproject.toml
  configs/
  data/
    tasks/
    preferences/
    splits/
  runs/
  experiments/
  reports/
  notebooks/
  scripts/
  src/agentalign/
    schemas.py
    tasks/
    verifiers/
    agents/
    tools/
    logging/
    preferences/
    training/
    evals/
    dashboard/
  tests/
```

## Naming Conventions

Use names that reveal role:

| Object     | Pattern           | Example                  |
|------------|-------------------|--------------------------|
| Task ID    | category_index    | py_fix_001               |
| Run ID     | date_agent_task   | 20260601_base_py_fix_001 |
| Experiment | date_method_model | exp_20260604_dpo_qwen15b |
| Dataset    | type_version      | preferences_v0.2.jsonl   |
| Report     | scope_date        | eval_dev_20260605.md     |

## Reproducibility

Minimum reproducibility packet:

- git commit SHA

- config file
- dataset version
- model name
- adapter path
- random seed
- GPU/runtime info
- exact command
- eval split
- metrics JSON

Your experiment is not reproducible if you cannot answer:

```
What command created this metric?
What data did it train on?
What code version ran it?
What model/adapters were loaded?
```

## Deterministic Seeds

Set seeds where possible:

```
import random
import numpy as np
import torch

def set_seed(seed: int) -> None:
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    if torch.cuda.is_available():
        torch.cuda.manual_seed_all(seed)
```

Be honest:

LLM generation and GPU kernels may still be nondeterministic. Record configs and run enough tasks to reduce noise.

## Dataset Versioning

Dataset versions should be immutable:

```
preferences_v0.1.jsonl
preferences_v0.1_report.md
preferences_v0.2.jsonl
preferences_v0.2_report.md
```

Dataset report should include:

- pair count
- task count
- task categories

- train/eval separation
- chosen score distribution
- rejected score distribution
- invalid pairs removed
- manual audit notes

## Experiment Tracking

For MVP, local experiment tracking is enough.

Do not start with a heavyweight platform. Use a consistent experiment folder. Add Weights & Biases, MLflow, Phoenix, or Langfuse only if it helps the demo.

Minimum `metrics.json`:

```
{
  "experiment_id": "exp_20260604_dpo_qwen15b",
  "base_pass_rate": 0.32,
  "tuned_pass_rate": 0.41,
  "base_invalid_action_rate": 0.18,
  "tuned_invalid_action_rate": 0.11,
  "eval_task_count": 50
}
```

## Clean CLI Design

Your CLI is the product interface for the repo.

Recommended commands:

```
agentalign tasks validate data/tasks
agentalign run --task-id py_fix_001 --agent baseline --out runs/dev
agentalign eval --config configs/eval_base.yaml
agentalign pairs build --runs runs/dev --out
data/preferences/preferences_v0.1.jsonl
agentalign train dpo --config configs/dpo_qwen15b.yaml
agentalign report eval --before runs/base --after runs/dpo
agentalign dashboard --runs runs
```

Good CLI properties:

- clear help text
- explicit config path
- fails loudly on missing files
- writes outputs to predictable folders
- prints the next useful command

## Error Handling Philosophy

Fail loudly at boundaries:

- task loading

- schema validation
- config parsing
- dataset loading
- experiment creation

Recover gracefully inside agent runs:

- invalid JSON
- tool timeout
- verifier failure
- model call failure

Never swallow errors without logging them.

## **Code Review Checklist**

Before calling a component done:

- Is there a typed schema?
- Is there a CLI path to run it?
- Does it save an artifact?
- Does it handle failure?
- Is there a small test?
- Are paths explicit?
- Is the output documented?
- Can another engineer understand it in 5 minutes?

## **How Startups Expect AI Engineers To Structure Repos**

Founders scan for:

- reproducible commands
- clean README
- small modules
- eval reports
- task/data examples
- clear metrics
- evidence of iteration
- failure analysis
- no dependency chaos

They do not expect:

- perfect infra
- Kubernetes
- huge model training
- paper-level novelty

They do expect:

Can this person take an ambiguous AI reliability problem and turn it into a working loop with data, evals, and evidence?

## 59. Research + Theory Map

The theory should serve the build. You are not trying to become a post-training researcher before writing code. You are learning enough theory to make better engineering decisions.

### Conceptual Dependency Graph

```
flowchart TD
    A["ML/DL basics"] --> B["Transformers and causal LM prompting"]
    B --> C["Tool-use agents"]
    C --> D["Trajectory logging"]
    D --> E["Verifier-based evaluation"]
    E --> F["Preference data"]
    F --> G["DPO"]
    G --> H["Before/after eval"]
    H --> I["Failure analysis"]
    I --> J["Iterative post-training"]

    E --> K["Observability"]
    K --> I
    C --> L["Sandboxing"]
    L --> E
```

### Exact Order Of Topics To Learn

| Order | Topic                       | Depth                         | Why                            |
|-------|-----------------------------|-------------------------------|--------------------------------|
| 1     | ReAct/tool-use agents       | Deep enough to implement loop | Agent trajectory format        |
| 2     | Verifier-based evals        | Deep                          | Core project credibility       |
| 3     | Pydantic/structured outputs | Deep                          | Reliable action parsing        |
| 4     | Preference datasets         | Deep                          | DPO input quality              |
| 5     | DPO                         | Deep enough to train/debug    | Main post-training method      |
| 6     | LoRA/QLoRA                  | Practical                     | Free GPU feasibility           |
| 7     | Observability/tracing       | Practical                     | Dashboard and failure analysis |



| Order | Topic               | Depth           | Why             |
|-------|---------------------|-----------------|-----------------|
| 8     | Docker/sandboxing   | Minimal for MVP | Safer execution |
| 9     | Advanced RL methods | Skim            | V2 vocabulary   |

## What To Deeply Study

Deeply study:

- ReAct agent loop structure
- DPO objective at implementation level
- TRL DPOTrainer inputs
- LoRA/QLoRA memory tradeoffs
- task/verifier methodology from SWE-bench and Terminal-Bench
- eval design and failure taxonomy
- structured schemas and validation

## What To Skim

Skim:

- full RLHF reward model pipeline
- PPO math
- GRPO details
- process supervision papers
- MCP protocol internals
- vLLM/SGLang production deployment
- OpenTelemetry implementation details

## What Theory Is Unnecessary For MVP

Skip for now:

- deriving policy gradients
- implementing PPO
- training reward models
- distributed inference serving
- Kubernetes
- multi-node RL
- full OpenTelemetry collector deployment

- production MCP server security architecture

These can be discussed conceptually in interviews, but they should not block the build.

## Minimal Theory Needed To Build This

You need to understand:

1. A language model maps prompt tokens to output tokens.
2. Tool-use agents convert model output into actions.
3. Verifiers turn messy behavior into measurable outcomes.
4. Preference learning uses comparisons, not absolute labels.
5. DPO tunes the model to prefer chosen responses over rejected responses.
6. LoRA updates a small set of adapter weights instead of full model weights.
7. Evaluation must use held-out tasks to avoid fooling yourself.
8. Observability means recording enough internal behavior to debug failures.

If you understand those eight things and build the loop, you have the project.

## Advanced Theory For Later

Learn after MVP:

- reward modeling
- process reward models
- step-level supervision
- online RL
- GRPO
- AI judges and judge calibration
- uncertainty and confidence intervals for evals
- constrained decoding internals
- distributed serving and batching

## RLHF vs DPO vs ORPO vs GRPO

| Method        | What it does                                       | Data needed                 | Complexity | MVP relevance           |
|---------------|----------------------------------------------------|-----------------------------|------------|-------------------------|
| RLHF with PPO | trains reward model, then optimizes policy with RL | preference labels + prompts | high       | understand conceptually |
| DPO           | directly optimizes chosen over rejected            | preference pairs            | medium     | use for MVP             |

| Method | What it does                                                                                          | Data needed           | Complexity | MVP relevance          |
|--------|-------------------------------------------------------------------------------------------------------|-----------------------|------------|------------------------|
| ORPO   | combines supervised learning and preference optimization without reference model in some formulations | preference-style data | medium     | mention as alternative |
| GRPO   | RL method often used for reasoning models, avoids separate value model compared with PPO-style setups | reward signals        | high       | V2/V3 only             |
| RLAIF  | preference/reward data generated by AI feedback                                                       | AI labels/verifiers   | medium     | project framing        |

The honest positioning:

AgentAlign uses verifier-derived preference data and DPO. It is not full RLHF, and that is a strength for a solo MVP because DPO is cheaper, simpler, and easier to evaluate.

## Modern Post-Training Landscape In 2026

The industry direction:

- more verifier-guided training for coding, math, and agents
- more synthetic preference data
- more tool-use and environment-based evals
- more process-level supervision, not only final answers
- more small specialized models and adapters
- more agent observability and trace-driven debugging
- more structured output and constrained decoding
- more emphasis on data quality than algorithm novelty for applied teams

Where AgentAlign fits:

It sits at the intersection of agent reliability, evals, post-training, and observability.

That is why it is internship-relevant.

## Theory-To-Code Translation Table

| Theory concept         | Code artifact            |
|------------------------|--------------------------|
| Agent trajectory       | Trajectory schema        |
| Tool action            | Action schema            |
| Reward/verifier signal | VerifierResult           |
| Preference pair        | PreferencePair JSONL     |
| DPO objective          | train_dpo.py             |
| Generalization         | held-out eval split      |
| Observability          | trace viewer/dashboard   |
| Failure taxonomy       | failure tags and reports |

## 60. Internship Positioning

This project should not be described as "I fine-tuned an LLM." That undersells it.

Describe it as:

I built an eval-driven post-training loop for terminal agents. The system creates verifiable coding/terminal tasks, runs an agent, logs trajectories, uses deterministic verifiers to generate preference pairs, fine-tunes a small model with DPO, and compares behavior before and after tuning through pass rate, invalid action rate, timeouts, and failure taxonomy.

## How Founders Evaluate Projects

Founders ask silently:

- Did this person build a real loop?
- Can they ship without handholding?
- Can they debug systems?
- Do they understand evals?
- Are the metrics credible?
- Can they explain tradeoffs?
- Did they use modern tools appropriately?
- Is the project small enough to be real?
- Did they learn from failures?

They are not impressed by:

- huge claims without evals
- "used LangChain" as the main achievement

- only a notebook
- no reproducibility
- no failure analysis
- vague "RLHF" wording without preference data

## What Recruiters Scan First

Recruiters scan:

1. README title
2. one-line project description
3. screenshots/GIF
4. architecture diagram
5. metrics table
6. tech stack
7. resume bullets

Make those readable in 30 seconds.

## "Tell Me About Your Project" Answer

Use this:

AgentAlign Lab is a small AI systems project focused on making terminal agents more reliable through verifier-guided preference learning. I built a pipeline where terminal tasks are executed in isolated workspaces, agent trajectories are logged step by step, deterministic verifiers score the outcome, and successful vs failed trajectories are converted into DPO preference pairs. Then I fine-tune a small code model with LoRA/QLoRA and evaluate whether behavior improves on held-out tasks.

The main engineering focus is not just training. It is the full loop: task design, tool execution, structured action parsing, verifier reliability, trajectory observability, preference dataset quality, DPO training, and failure analysis. The metrics I track include pass rate, invalid action rate, timeout rate, average steps, and failure-tag distribution before and after tuning.

Shorter version:

I built a verifier-driven post-training and eval pipeline for terminal agents: run tasks, log traces, score with deterministic verifiers, generate preference pairs, DPO-tune a small model, and analyze reliability improvements.

## How To Explain Failures Honestly

Good answer:

The first DPO run did not improve pass rate. I traced it to noisy preference pairs: many chosen and rejected trajectories were both failures with only small score differences. I added a score-margin filter, audited random pairs, and separated failure tags. After that, the data became cleaner and the invalid action rate improved, even when pass rate improved only modestly.

Bad answer:

The model was bad, so I tried another model.

Founders like honest iteration. They dislike vague blaming.

## How To Discuss Tradeoffs

Use this pattern:

I chose X instead of Y because the project goal was Z. The tradeoff was A, but it let me ship B.

Examples:

- "I used deterministic verifiers instead of LLM judges because the MVP needed credible labels. The tradeoff is narrower task types, but the evals are more trustworthy."
- "I used a custom agent loop instead of a framework because I needed full control over traces and parser failures."
- "I deferred Docker because I was new to infra and local tempdir execution was enough for MVP. I documented Docker as V2 safety work."
- "I used Qwen2.5-Coder-1.5B because it fits free GPU constraints and still supports coding behavior."

## How To Present Metrics Credibly

Bad metric:

Accuracy improved by 30 percent.

Good metric:

On 50 held-out terminal tasks, pass rate improved from 32 percent to 41 percent, invalid action rate dropped from 18 percent to 11 percent, and timeout rate stayed roughly flat. The biggest remaining failure mode was wrong-file edits.

Always include:

- denominator
- split name
- before/after
- metric definition
- limitations

## Expected Interview Questions

| Question                                 | What they test                   |
|------------------------------------------|----------------------------------|
| Why DPO instead of PPO?                  | practical post-training judgment |
| How do you know the verifier is correct? | eval credibility                 |

| Question                              | What they test        |
|---------------------------------------|-----------------------|
| How did you avoid train/eval leakage? | scientific discipline |
| What was the hardest bug?             | debugging depth       |
| What did DPO improve?                 | evidence over hype    |
| What failed?                          | honesty and iteration |
| Why terminal agents?                  | startup relevance     |
| How would you scale this?             | systems thinking      |
| Why not use LangGraph?                | tool judgment         |
| How would you add observability?      | infra awareness       |

## Hard Technical Interview Questions And Answers

Question: Why is verifier quality central?

Answer:

Because the verifier creates the supervision signal. If it has false positives, the system may train on bad behavior as chosen. If it has false negatives, it may reject good behavior. I treated verifier development like eval engineering: reference solutions must pass, broken templates must fail, timeouts are structured failures, and failure tags are logged for analysis.

Question: Why DPO instead of supervised fine-tuning?

Answer:

SFT can teach a single desired behavior, but DPO uses contrastive feedback: for the same task prompt, one trajectory is preferred over another. Since terminal agents fail in many ways, like invalid actions or looping, chosen/rejected pairs are a natural way to teach reliability preferences.

Question: How do you know improvement is not overfitting?

Answer:

I split by task ID before generating training pairs, froze a held-out eval split, and reported before/after metrics on that split. I also track multiple metrics, not just pass rate, because a model can overfit pass rate while becoming worse at action formatting or timeouts.

Question: What is the main risk in verifier-guided DPO?

Answer:

The main risk is learning verifier artifacts instead of robust behavior. For example, if tests are weak, the model may learn shortcuts. That is why I include manual pair audits, failure tags, and reference-solution checks.

Question: What would you improve with more time?

Answer:

I would add Docker sandboxing, step-level rewards, AI-assisted failure tagging, and trace export to Phoenix or Langfuse. I would also expand from final-outcome

preferences to process preferences, because many terminal-agent failures happen before the final verifier result.

## Founder-Style Grilling Simulation

Founder: "Why should I care? Everyone has an agent demo."

Answer:

This is not only an agent demo. It is an agent reliability loop. The important part is that every run creates data: traces, verifier scores, failure tags, and preference pairs. That lets the system improve and lets engineers understand why it fails.

Founder: "Are your metrics real?"

Answer:

They are small-scale but real. I use held-out tasks, deterministic verifiers, and before/after comparisons. I also show failure breakdowns and limitations instead of claiming broad generalization.

Founder: "Why did you not just use an agent framework?"

Answer:

I studied LangGraph and the OpenAI Agents SDK, but for the MVP I wanted explicit control over parser failures, tool execution, trace schema, and preference pair generation. Frameworks are useful later; transparency mattered more early.

Founder: "What was the hardest engineering part?"

Answer:

The hardest part was making the data trustworthy. It is easy to run agents and collect messy traces. It is harder to ensure tasks are valid, verifiers are not lying, pairs have meaningful contrast, and eval results are not inflated.

## How To Sound Like An AI Systems Engineer Instead Of A Tutorial Follower

Say:

- "I used eval-driven development."
- "I separated model behavior from verifier quality."
- "I tracked invalid action rate because pass rate alone hides parser failures."
- "I split by task ID to avoid leakage."
- "I used deterministic verifiers first and deferred LLM judges."
- "I treated preference data as the main bottleneck."
- "I logged full trajectories so I could debug failures."

Avoid:

- "I used RLHF to make an agent better" if you did DPO.



- "The model learned" without metrics.
  - "I followed a tutorial."
  - "LangChain did the agent part."
  - "Accuracy improved" without denominator.
- 

## 61. README + Demo Masterclass

The README is your recruiter interface. The demo is your founder interface.

### Exact README Structure

Use this structure:

```
# AgentAlign Lab

One-line pitch.

![demo gif](assets/demo.gif)

## What It Does
Short explanation of the loop.

## Why It Matters
Agent reliability, evals, post-training, observability.

## Architecture
Mermaid or image diagram.

## Results
Before/after metrics table.

## Demo
Commands + dashboard screenshot.

## Core Features
- terminal task runner
- verifier system
- trajectory logging
- preference pair generation
- DPO training
- eval/failure dashboard

## How To Run
Commands.

## Dataset
Task types and pair quality.

## Evaluation Methodology
Held-out split, metrics, limitations.

## Failure Analysis
Top failure modes.

## Tech Stack
Models, TRL, PEFT, Gradio, Pydantic, etc.
```

## Roadmap  
MVP done, V2 ideas.

## Best Screenshot Strategy

Include 4 images:

1. architecture diagram
2. dashboard metrics page
3. trace viewer showing one failed and one successful run
4. before/after eval table

Do not use only terminal screenshots. Terminal screenshots are useful, but founders need to see system-level clarity.

## Best GIF Strategy

One GIF should show:

```
select task -> run agent -> see trajectory -> verifier result -> dashboard update
```

Keep it under 30 seconds.

Do not show:

- long model waiting
- scrolling logs for 2 minutes
- only a static landing page

## Architecture Diagram Guidance

Your architecture diagram should show data movement:

```
flowchart LR
    A["Task suite"] --> B["Agent runner"]
    B --> C["Tool executor"]
    C --> D["Trajectory logs"]
    D --> E["Verifier"]
    E --> F["Preference pairs"]
    F --> G["DPO training"]
    G --> H["Tuned agent"]
    H --> I["Evaluation"]
    I --> J["Dashboard"]
```

Use labels founders understand:

- "Verifier"
- "Trajectory logs"
- "Preference pairs"
- "DPO fine-tuning"

- "Held-out evaluation"

Avoid labels only researchers understand:

- "policy"
- "KL"
- "implicit reward"

Use those in technical sections, not the hero diagram.

## Metrics Presentation

README metrics table:

| Metric              | Base agent | DPO-tuned agent | Change |
|---------------------|------------|-----------------|--------|
| Pass rate           | 32%        | 41%             | +9 pp  |
| Invalid action rate | 18%        | 11%             | -7 pp  |
| Timeout rate        | 14%        | 12%             | -2 pp  |
| Avg steps           | 6.8        | 6.1             | -0.7   |

Add caveat:

Measured on 50 held-out tasks. Small-scale student project; numbers are intended to show methodology and relative improvement, not broad benchmark claims.

That caveat makes you look more credible, not weaker.

## Failure Analysis Presentation

Show failure modes:

| Failure tag    | Base count | Tuned count | Interpretation              |
|----------------|------------|-------------|-----------------------------|
| invalid_action | 9          | 5           | DPO helped schema following |
| looping        | 6          | 4           | modest improvement          |
| wrong_file     | 5          | 7           | still unsolved              |
| timeout        | 7          | 6           | mostly unchanged            |

This tells a stronger story than only showing wins.

## Dashboard Storytelling

The dashboard should answer:

1. Is the tuned model better?
2. On which task types?
3. Which failure modes improved?

4. Which got worse?
5. Can I inspect one trace?
6. Can I inspect one preference pair?

## **What Founders Care About Visually**

They care about:

- clear architecture
- evidence of working system
- metrics with denominators
- failure breakdown
- trace inspectability
- clean repo
- fast demo path

They do not care about:

- decorative UI
- huge hero sections
- abstract AI art
- overdesigned landing page

## **How To Make This Project Look Elite On GitHub**

Checklist:

- repo has a clear one-line pitch
- README has a GIF above the fold
- architecture diagram appears in first screen
- results table is visible without scrolling too far
- demo command works
- sample data included
- reports folder has eval writeups
- code is modular
- issues/roadmap show thoughtful next steps
- failures are documented honestly
- dependencies are pinned enough to run

Add these files:

```
README.md
docs/architecture.md
docs/evaluation.md
docs/failure_analysis.md
docs/training.md
reports/base_vs_dpo_eval.md
assets/demo.gif
assets/dashboard.png
```

Demo Script

Use this 90-second demo:

- 1. "This is AgentAlign Lab, a verifier-guided post-training loop for terminal agents."
- 2. Show architecture diagram.
- 3. Run one task through the base agent.
- 4. Show trajectory and verifier result.
- 5. Show preference pair generation.
- 6. Show before/after DPO eval table.
- 7. Show failure dashboard.
- 8. End with what improved, what failed, and what you would build next.

This sounds like engineering, not hype.

62. Modern AI Infrastructure Awareness (2026)

You do not need to deeply implement every modern tool. But you should know what each tool is, why it exists, and whether it belongs in AgentAlign now or later.

Tool Awareness Matrix

| Tool              | What it is                                                      | Learn deeply?            | Use in MVP?        | Why                                           |
|-------------------|-----------------------------------------------------------------|--------------------------|--------------------|-----------------------------------------------|
| vLLM              | high-throughput LLM serving engine                              | conceptually + basic use | no                 | useful for serving, overkill for Colab MVP    |
| SGLang            | high-performance serving/runtime for LLMs and multimodal models | conceptually             | no                 | important for RL/serving infra, too heavy now |
| Ollama            | local model runner                                              | basic                    | optional           | useful local inference, not training          |
| LangGraph         | graph/state-machine framework for agents                        | medium later             | no                 | good reference, custom loop better for MVP    |
| OpenAI Agents SDK | agent runtime with tools, guardrails, tracing, sandbox          | conceptually             | no core dependency | excellent design reference                    |

| Tool                 | What it is                                        | Learn deeply?     | Use in MVP?      | Why                                            |
|----------------------|---------------------------------------------------|-------------------|------------------|------------------------------------------------|
|                      | agents                                            |                   |                  |                                                |
| MCP                  | protocol for connecting agents to tools/resources | recognize         | V3               | important standard, not needed for local tools |
| Phoenix              | open-source observability/evals                   | practical later   | V2 optional      | great for traces/evals                         |
| Langfuse             | LLM observability, prompts, evals                 | practical later   | V2 optional      | startup-relevant observability                 |
| Braintrust           | evals and experiment platform                     | conceptually      | no               | learn eval workflow ideas                      |
| DSPy                 | framework for programming/optimizing LM pipelines | conceptually      | no               | useful mindset, not core                       |
| Structured outputs   | schema-constrained model outputs                  | deeply            | yes conceptually | action reliability                             |
| Constrained decoding | decoding constrained by JSON/schema/grammar       | conceptually      | optional         | helps reduce invalid JSON                      |
| OpenTelemetry        | tracing standard                                  | recognize         | no               | V2/V3 observability                            |
| Eval platforms       | systematic testing systems                        | deeply in concept | local MVP        | core project theme                             |

## vLLM

Official docs: <https://docs.vllm.ai/>

What it is:

vLLM is a high-throughput serving engine for LLM inference. It matters when you want fast, batched, OpenAI-compatible serving of models.

Why startups use it:

- high throughput
- OpenAI-compatible server
- batching
- quantization support

- LoRA serving
- structured outputs
- observability hooks

For AgentAlign:

- Do not use for MVP.
- Mention it in V2 as a serving option.
- Study structured outputs and OpenAI-compatible serving conceptually.

Best learning target:

Understand why serving is separate from training and why batching/latency matter.

## SGLang

Official docs: <https://docs.sglang.io/>

What it is:

SGLang is a production-oriented serving/runtime framework for large language and multimodal models, focused on low latency and high throughput.

Why startups care:

- inference performance
- prefix caching
- OpenAI-compatible APIs
- RL/post-training integrations
- large-scale serving

For AgentAlign:

- Do not use in MVP.
- Mention as modern serving/RL infra awareness.
- Use only if you later deploy a larger evaluation service.

## Ollama

Official repo/docs: <https://github.com/ollama/ollama>

What it is:

Ollama makes local model running easy.

For AgentAlign:

- useful for quick local demos on CPU/Mac if model quality is enough
- not a training tool

- not ideal for DPO

Learn:

- how local model serving differs from HF Transformers training
- how OpenAI-compatible/local APIs simplify model clients

## LangGraph

Official docs: <https://docs.langchain.com/oss/python/langgraph/workflows-agents>

What it is:

LangGraph helps build stateful agent workflows as graphs.

For AgentAlign:

- Study for architecture ideas.
- Do not use in MVP unless your custom loop becomes unmanageable.
- It is more useful when your agent has multiple branches, memory states, or human-in-loop paths.

Correct positioning:

I used a custom loop for trace/control in MVP, while studying LangGraph as a production workflow reference.

## OpenAI Agents SDK

Official docs: <https://openai.github.io/openai-agents-python/>

What it is:

An agent SDK with agents, tools, guardrails, handoffs, sessions, tracing, and sandbox-agent concepts.

For AgentAlign:

- Excellent design reference.
- Do not make it a hard dependency for MVP.
- Study how it structures tools, tracing, and guardrails.

Useful comparison point:

AgentAlign reimplements a minimal transparent agent loop to expose exactly the data needed for verifier-guided preference learning.

## MCP

Official latest spec: <https://modelcontextprotocol.io/specification/latest>

What it is:

MCP standardizes how LLM applications connect to external tools, resources, and prompts.



For AgentAlign:

- Recognize it.
- Do not implement it in MVP.
- V3 idea: expose terminal-task tools as MCP tools.

Important caution:

MCP expands tool access. Tool access expands risk. Learn the security model before exposing real filesystem or shell capabilities.

## Phoenix

Official docs: <https://arize.com/docs/phoenix>

What it is:

Phoenix is open-source AI observability and evaluation tooling. It supports traces, evals, datasets, and experiments.

For AgentAlign:

- Strong V2 option.
- Use if you want to export traces to a real observability system.
- MVP can mimic the concepts locally with JSONL + dashboard.

## Langfuse

Official docs: <https://langfuse.com/docs>

What it is:

Open-source LLM engineering platform for observability, prompt management, evaluation, and experiments.

For AgentAlign:

- Strong V2 option.
- Excellent to mention for evals/observability startup relevance.
- Do not integrate before local dashboard works.

## Braintrust

Official docs: <https://www.braintrust.dev/docs/evaluate>

What it is:

An eval and observability platform focused on systematic evaluation and iteration.

For AgentAlign:

- Study the eval workflow.
- Do local eval implementation first.

- Mention Braintrust concepts in README resources if relevant.

## DSPy

Official docs: <https://dspy.ai/>

What it is:

DSPy is a framework for programming and optimizing language-model pipelines as modules rather than prompt strings.

For AgentAlign:

- Useful conceptual reference for eval-driven LM programming.
- Not needed for MVP.
- Potential V2: optimize prompts for task generation or failure tagging.

## Structured Outputs And Constrained Decoding

Useful resources:

- OpenAI structured outputs guide: <https://platform.openai.com/docs/guides/structured-outputs>
- vLLM structured outputs: [https://docs.vllm.ai/en/latest/features/structured\\_outputs/](https://docs.vllm.ai/en/latest/features/structured_outputs/)
- Pydantic models: <https://docs.pydantic.dev/latest/concepts/models/>

Why it matters:

Agents need machine-readable actions. Free-form text is fragile.

MVP approach:

- prompt for strict JSON
- validate with Pydantic
- retry once
- log invalid actions

V2 approach:

- use provider structured outputs where available
- use constrained decoding if running served models
- compare invalid action rate before/after

## Which Tools Are Hype For This MVP?

Overkill now:

- Kubernetes
- distributed vLLM/SGLang serving
- full MCP server

- online RL infrastructure
- multi-agent orchestration framework
- OpenTelemetry collector stack
- reward model training

Useful now:

- Pydantic
- TRL
- PEFT/QLoRA
- Hugging Face Datasets
- Gradio
- pytest
- ruff
- JSONL traces
- local eval reports

Deeply learn:

- DPO/TRL
- verifiers/evals
- structured parsing
- failure analysis
- preference data quality

Recognize conceptually:

- vLLM
- SGLang
- MCP
- OpenAI Agents SDK
- LangGraph
- Phoenix/Langfuse/Braintrust
- DSPy

---

## 63. System Design Visuals Addendum

Use these diagrams in README, docs, presentation, and interviews. They make the system legible fast.

## High-Level Architecture

```
flowchart TD
    A["Task suite"] --> B["Agent runner"]
    B --> C["Model client"]
    B --> D["Action parser"]
    B --> E["Tool executor"]
    E --> F["Workspace"]
    B --> G["Trajectory logger"]
    G --> H["Verifier"]
    H --> I["Evaluation metrics"]
    G --> J["Preference pair builder"]
    J --> K["DPO dataset"]
    K --> L["DPO trainer"]
    L --> M["Tuned adapter"]
    M --> N["Tuned agent eval"]
    I --> O["Dashboard"]
    N --> O
```

## Artifact Flow

```
flowchart LR
    A["task.json"] --> B["workspace copy"]
    B --> C["trajectory.json"]
    C --> D["verifier_result.json"]
    C --> E["preference_pair.jsonl"]
    E --> F["dpo_dataset_vX"]
    F --> G["lora_adapter"]
    G --> H["eval_results.json"]
    H --> I["report.md"]
    H --> J["dashboard"]
```

## Agent Run Sequence

```
sequenceDiagram
    participant User
    participant CLI
    participant Runner
    participant Model
    participant Parser
    participant Tool
    participant Verifier
    participant Logger

    User->>CLI: agentalign run --task-id py_fix_001
    CLI->>Runner: load task and config
    Runner->>Logger: create trajectory
    loop max_steps
        Runner->>Model: prompt with task + history
        Model-->>Runner: raw action text
        Runner->>Parser: parse and validate
        Parser-->>Runner: Action or error
        Runner->>Tool: execute valid action
        Tool-->>Runner: observation
        Runner->>Logger: append step
    end
    Runner->>Verifier: run verifier
    Verifier-->>Runner: VerifierResult
    Runner->>Logger: finalize trajectory
```

## Verifier Pipeline

```
flowchart TD
    A["Final workspace"] --> B["Pre-check protected files"]
    B --> C["Run verifier command"]
    C --> D{"Exit code 0?"}
    D -- "yes" --> E["passed=true score=1.0"]
    D -- "no" --> F["passed=false score=0.0"]
    E --> G["Capture stdout/stderr"]
    F --> G
    G --> H["Assign failure tags"]
    H --> I["Validate VerifierResult"]
    I --> J["Persist result"]
```

## Preference Dataset Pipeline

```
flowchart TD
    A["Trajectories"] --> B["Group by task"]
    B --> C["Filter valid verifier results"]
    C --> D["Rank by score"]
    D --> E{"Score margin enough?"}
    E -- "no" --> F["Skip ambiguous pair"]
    E -- "yes" --> G["Create chosen/rejected pair"]
    G --> H["Validate schema"]
    H --> I["Audit sample"]
    I --> J["Write preference JSONL"]
    J --> K["Dataset report"]
```

## DPO Workflow

```
flowchart LR
    A["Base code model"] --> B["Load tokenizer"]
    C["Preference JSONL"] --> D["Format prompt/chosen/rejected"]
    D --> E["Tokenize"]
    A --> F["LoRA/QLoRA config"]
    E --> G["DPOTrainer"]
    F --> G
    G --> H["Adapter checkpoint"]
    H --> I["Eval tuned agent"]
    I --> J["Compare with base"]
```

## Evaluation Pipeline

```
flowchart TD
    A["Frozen eval split"] --> B["Run base agent"]
    A --> C["Run tuned agent"]
    B --> D["Base trajectories"]
    C --> E["Tuned trajectories"]
    D --> F["Aggregate metrics"]
    E --> F
    F --> G["Pass rate"]
    F --> H["Invalid action rate"]
    F --> I["Timeout rate"]
    F --> J["Failure tags"]
    G --> K["Eval report"]
    H --> K
    I --> K
    J --> K
```

## Trace Lifecycle

```
stateDiagram-v2
    [*] --> Created
    Created --> Running
    Running --> StepLogged
    StepLogged --> Running
    Running --> Verification
    Verification --> Finalized
    Verification --> FailedButSaved
    FailedButSaved --> Finalized
    Finalized --> Dashboard
    Finalized --> PreferenceBuilder
    Finalized --> FailureAnalysis
```

## Dashboard Data Flow

```
flowchart TD
    A["runs/*.json"] --> B["Load trajectories"]
    C["reports/*.json"] --> B
    B --> D["Compute aggregate metrics"]
    B --> E["Build failure table"]
    B --> F["Build trace index"]
    D --> G["Metric cards"]
    E --> H["Failure charts"]
    F --> I["Trace viewer"]
    G --> J["Gradio dashboard"]
    H --> J
    I --> J
```

## Low-Level Module Boundary Diagram

```
flowchart TB
    subgraph Schemas
        S1["Task"]
        S2["Action"]
        S3["Trajectory"]
        S4["VerifierResult"]
        S5["PreferencePair"]
    end

    subgraph Runtime
        R1["AgentRunner"]
        R2["ModelClient"]
        R3["ActionParser"]
        R4["ToolExecutor"]
        R5["TrajectoryLogger"]
    end

    subgraph Learning
        L1["PairBuilder"]
        L2["DPOTrainer script"]
        L3["Evaluator"]
    end

    Schemas --> Runtime
    Runtime --> Learning
```

## 64. Execution Realism

This project is very buildable. It is also very easy to ruin through scope creep.

### Why Students Fail Projects Like This

Common failure pattern:

```
read too much -> choose too many tools -> build half of everything -> no working loop -> panic -> polish README with weak results
```

Specific reasons students fail:

- they start with DPO before tasks/verifiers work
- they chase novelty instead of reliability
- they use a huge framework and cannot debug it
- they build dashboard before data exists
- they train without a frozen eval split
- they report metrics they cannot reproduce
- they avoid ugly debugging logs
- they keep changing the project name/scope

### How Scope Creep Kills This Project

Scope creep often disguises itself as ambition:

- "I should add browser agents."
- "I should implement MCP."
- "I should use vLLM."
- "I should train a reward model."
- "I should make a full benchmark."
- "I should add multi-agent debate."

Those may be good ideas later. In weeks 1-4, they are mostly project killers.

The MVP has one sacred loop:

```
task -> agent -> trajectory -> verifier -> preference pair -> DPO -> eval -> dashboard
```

Anything that does not strengthen that loop is optional.

### When To Cut Features

Cut a feature if:

- it does not affect demo or metrics
- it delays the first full loop

- it requires learning a new tool for more than one day
- it adds infra risk without improving evidence
- you cannot explain why a founder would care
- it is only there because it sounds advanced

Keep a feature if:

- it makes evals more credible
- it improves trace visibility
- it produces better preference data
- it reduces common failures
- it makes the demo clearer
- it strengthens interview stories

## When To Stop Researching And Start Coding

Start coding when you can write:

Input:  
Output:  
Success condition:  
Failure condition:

Example:

Input: task JSON and workspace template.  
Output: verifier\_result.json.  
Success: broken workspace fails, reference solution passes.  
Failure: verifier cannot run or gives wrong pass/fail.

If you can write that, you can build.

## How To Know If A Feature Is Worth Building

Score it:

| Question                               | Score 0-2 |
|----------------------------------------|-----------|
| Does it improve the demo?              |           |
| Does it improve metrics credibility?   |           |
| Does it improve learning signal?       |           |
| Can it be built in less than two days? |           |
| Can it become an interview story?      |           |

Build if score  $\geq 7$ .

Defer if score  $\leq 5$ .



# Things That Feel Productive But Are Actually Avoidance

- changing UI colors before metrics exist
- reorganizing folders repeatedly
- reading every DPO blog
- watching agent framework tutorials
- comparing 10 base models
- making a logo
- writing a long paper before code works
- adding Docker before local runner works
- adding Phoenix before local traces work
- trying to make tasks "benchmark-grade" before you have 10 working tasks

## High-Leverage Work vs Fake Work

| High-leverage work         | Fake work                       |
|----------------------------|---------------------------------|
| add 10 verified tasks      | browse benchmark papers all day |
| fix parser failures        | redesign prompt endlessly       |
| audit preference pairs     | generate 1000 noisy pairs       |
| run held-out eval          | show cherry-picked demos        |
| document a failed DPO run  | hide bad results                |
| build trace viewer         | build decorative landing page   |
| write reproducible command | rely on notebook state          |

## 65. AI Engineering Mindset

AI systems engineers think in loops, not demos.

### How AI Systems Engineers Think

They ask:

- What is the input distribution?
- What is the failure mode?
- How do we measure it?
- What artifact proves it?
- What changed between runs?
- Can we reproduce it?

- Is the model problem actually a data/eval/tooling problem?
- What is the smallest intervention?

## Debugging Mindset

Do not debug by vibes. Debug by layers:

schema -> data -> runner -> tool -> verifier -> model -> training -> eval

Do not ask "why is the model bad?"

Ask:

Which layer produced the first wrong artifact?

## Eval-Driven Development

Every important change should connect to an eval:

- parser change -> invalid action rate
- verifier change -> false positive/negative checks
- task change -> task validation report
- DPO change -> held-out pass rate
- prompt change -> before/after trajectory comparison
- dashboard change -> faster failure diagnosis

## Iterative Engineering

The right order:

make it run -> make it visible -> make it correct -> make it better -> make it pretty

Not:

make it impressive -> make it complex -> hope it works

## Startup Engineering Mentality

A startup does not need you to know everything. It needs you to:

- learn fast
- isolate problems
- ship narrow loops
- communicate tradeoffs
- avoid drama around bugs
- produce evidence
- move from ambiguity to artifact

This project can show exactly that.

## Measuring Progress

Progress is not hours spent.

Progress is:

- tasks added and validated
- verifiers proven
- trajectories collected
- pairs audited
- eval runs completed
- failure modes reduced
- README improved with evidence

## Avoiding Perfectionism

Perfectionism says:

I need the ideal architecture before starting.

Engineering says:

I need the smallest architecture that can produce truthful feedback.

## Learning Through Building

The project itself is the course.

Each component teaches:

| Component        | Skill learned           |
|------------------|-------------------------|
| Task suite       | eval design             |
| Verifier         | reliability engineering |
| Agent loop       | tool-use systems        |
| Parser           | structured outputs      |
| Trajectory logs  | observability           |
| Preference pairs | alignment data          |
| DPO              | post-training           |
| Dashboard        | failure analytics       |
| README           | technical communication |

## 66. Advanced Optional V2/V3 Ideas

These are optional expansions. They are not MVP requirements.

### V2/V3 Idea Matrix

| Idea                      | Realistic? | Startup signal | Risk            | Recommendation              |
|---------------------------|------------|----------------|-----------------|-----------------------------|
| Docker sandbox            | high       | high           | medium          | best V2                     |
| Phoenix/Langfuse tracing  | high       | high           | medium          | strong V2                   |
| AI failure tagger         | high       | medium         | low             | useful V2                   |
| Step-level rewards        | medium     | high           | medium          | V2/V3                       |
| Process supervision       | medium     | high           | high            | V3                          |
| Synthetic task generation | medium     | medium         | high data noise | only after task validator   |
| Online RL                 | low        | high           | very high       | discuss, do not build early |
| Reward model              | low-medium | medium         | high            | defer                       |
| GRPO experiments          | low        | high           | high            | V3 only                     |
| Browser agents            | medium     | medium         | high scope      | separate project            |
| MCP tools                 | medium     | medium         | security risk   | V3                          |
| Distributed eval          | low        | medium         | overkill        | skip                        |
| Multi-agent workflows     | medium     | medium         | scope creep     | only if core loop done      |

### Online RL

What it means:

The model improves through interaction, not only offline preference pairs.

Why it is attractive:

- closer to frontier post-training
- useful for agents
- can optimize against verifiers

Why it is too much now:

- unstable
- expensive
- harder to debug

- needs careful reward design

Recommendation:

Mention as future work. Do not build in MVP.

## AI Judges

Use when:

- tasks are semantic
- deterministic verifier is impossible
- you need qualitative scoring

Risks:

- judge bias
- inconsistency
- cost
- prompt sensitivity

Best V2 use:

AI judge assists failure tagging, not final pass/fail.

## Step-Level Rewards And Process Supervision

Why useful:

Terminal agents often fail before final output. Step-level signals can teach better intermediate behavior.

Example:

Good step: inspect file before editing.  
Bad step: edit random file without reading.

V2 approach:

- tag trajectories at step level
- create chosen/rejected step snippets
- train on process-style preferences

## Synthetic Data Generation

Useful for:

- generating candidate tasks
- generating buggy files
- generating rejected trajectories

Danger:

Synthetic data can be confidently wrong. Every generated task still needs deterministic validation.

Rule:

Synthetic tasks are candidates, not dataset entries.

## Self-Play Improvements

Idea:

Use multiple agent attempts per task and select winners via verifier.

Realistic V2:

- sample 3-5 attempts per task
- verifier ranks them
- generate pairs automatically

This is a strong extension because it scales preference data without human labels.

## MCP Tools

V3 idea:

Expose task operations as MCP tools:

- read file
- write file
- run tests
- inspect workspace

Why interesting:

It connects your project to the modern agent-tool ecosystem.

Why risky:

Tool permissions and filesystem access become security-sensitive.

## OpenTelemetry Tracing

V2/V3:

- export local trace spans to Phoenix or Langfuse
- map agent step to span
- map tool call to span
- map verifier call to span

Why startups care:

This shows you understand agent observability beyond print logs.

## Reward Models

V3 only.

Why:

Training a reward model adds complexity and may not beat deterministic verifiers for your task type.

Better story:

For MVP, deterministic verifiers provide cheaper and more reliable supervision. Reward modeling is future work for tasks where deterministic verification is impossible.

## GRPO Experiments

V3 only.

Use as conceptual awareness, not MVP implementation. It is useful to know that modern reasoning/post-training work increasingly explores RL-style optimization, but building GRPO correctly is not necessary for internship signal here.

---

## 67. Implementation Roadmap 2.0

This roadmap converts the 6-8 week plan into deliverables, risk checkpoints, fallback strategies, and GitHub milestones.

The north-star rule:

By the end of Week 3, you must have the full loop working on a tiny dataset, even if the model improvement is not impressive yet.

### Week 1: Task + Verifier Foundation

Goal:

Create verifiable terminal tasks and run deterministic verifiers in isolated workspaces.

Daily deliverables:

|  | Day | Deliverable             | Success condition                                 |
|--|-----|-------------------------|---------------------------------------------------|
|  | 1   | repo skeleton + schemas | Task, Action, Trajectory, VerifierResult validate |
|  | 2   | 5 handwritten tasks     | broken templates fail verifier                    |
|  | 3   | reference fixes         | reference solutions pass verifier                 |
|  | 4   | workspace runner        | task copies to temp workspace                     |
|  | 5   | verifier CLI            | agentalign verify --task-id X works               |

|  | Day | Deliverable            | Success condition                    |
|--|-----|------------------------|--------------------------------------|
|  | 6   | task validation report | report lists pass/fail sanity checks |
|  | 7   | cleanup/docs           | README shows task/verifier flow      |

Exact success condition:

I can run 5 tasks through verifiers, prove broken versions fail, prove reference fixes pass, and save structured verifier results.

Danger zone:

- reading agent papers instead of writing tasks
- writing tasks without reference solutions
- verifier passes broken code

Fallback:

- reduce to 3 tasks
- use only Python/pytest
- skip agent entirely until verifier works

GitHub milestone:

v0.1-task-verifier-foundation

## Week 2: Agent Loop + Tool Execution

Goal:

Run a baseline terminal agent and log every step.

Daily deliverables:

|  | Day | Deliverable                | Success condition                   |
|--|-----|----------------------------|-------------------------------------|
|  | 8   | model client stub          | fake model can drive loop           |
|  | 9   | action parser              | valid/invalid JSON tests pass       |
|  | 10  | tool executor              | read/write/run command actions work |
|  | 11  | baseline model integration | one real model attempt saved        |
|  | 12  | trajectory logger          | every step saved as JSON            |
|  | 13  | max steps/timeouts         | looping cannot hang run             |
|  | 14  | 10-task baseline run       | report generated                    |

Exact success condition:



The agent can attempt 10 tasks, save trajectories, run verifier, and produce pass/fail metrics.

Danger zone:

- model output parsing dominates everything
- agent loop becomes one giant function
- no failed runs saved

Fallback:

- use scripted/random baseline first
- reduce action schema to `read_file`, `write_file`, `run_command`, `stop`
- allow only one parser retry

GitHub milestone:

v0.2-agent-runner-traces

## Week 3: Preference Pair Pipeline

Goal:

Generate clean chosen/rejected pairs from verifier-ranked trajectories.

Daily deliverables:

|  | Day | Deliverable                    | Success condition                       |
|--|-----|--------------------------------|-----------------------------------------|
|  | 15  | run multiple attempts per task | at least 2 trajectories/task            |
|  | 16  | pair schema                    | PreferencePair validates                |
|  | 17  | pair builder                   | same-task chosen/rejected pairs created |
|  | 18  | pair filters                   | score margin + failure tag filters      |
|  | 19  | manual audit                   | 20 pairs audited with notes             |
|  | 20  | dataset report                 | stats for pair count/tasks/tags         |
|  | 21  | freeze v0.1 dataset            | JSONL and report committed              |

Exact success condition:

I have 300+ validated preference pairs with a dataset report and manual audit notes.

Danger zone:

- pairs are noisy
- chosen/rejected labels swapped
- same task appears in train and eval

Fallback:

- make fewer but cleaner pairs for the first DPO smoke test
- manually curate 50 high-confidence pairs before scaling
- delay DPO until pair quality is credible

GitHub milestone:

v0.3-preference-dataset

## Week 4: First DPO Run

Goal:

Train a small LoRA/QLoRA adapter and compare before/after behavior.

Daily deliverables:

|  | Day | Deliverable                | Success condition             |
|--|-----|----------------------------|-------------------------------|
|  | 22  | DPO data formatting        | TRL dataset loads             |
|  | 23  | 10-pair overfit smoke test | training runs without crash   |
|  | 24  | full small DPO run         | adapter saved                 |
|  | 25  | adapter loading            | tuned model generates actions |
|  | 26  | before/after eval          | same split, same config       |
|  | 27  | experiment report          | metrics + samples             |
|  | 28  | debugging pass             | fix biggest obvious issue     |

Exact success condition:

A DPO adapter trains, loads, and is evaluated against the base model on held-out tasks.

Danger zone:

- CUDA OOM
- no held-out eval
- adapter not actually loaded
- tuned model breaks JSON

Fallback:

- reduce model size
- reduce sequence length
- reduce pairs
- train one epoch
- report negative result honestly if needed

GitHub milestone:

v0.4-first-dpo-eval

## Week 5: Failure Analysis + Dashboard

Goal:

Turn raw runs into visible reliability analysis.

Daily deliverables:

|  | Day | Deliverable            | Success condition                   |
|--|-----|------------------------|-------------------------------------|
|  | 29  | failure taxonomy       | tags applied to failed runs         |
|  | 30  | metrics aggregator     | pass/invalid/timeout/steps computed |
|  | 31  | trace viewer           | inspect one run step by step        |
|  | 32  | before/after dashboard | base vs tuned comparison visible    |
|  | 33  | pair inspector         | chosen/rejected examples visible    |
|  | 34  | failure report         | top 5 failure modes explained       |
|  | 35  | dashboard screenshots  | assets ready for README             |

Exact success condition:

The dashboard can answer what improved, what failed, and show the trace evidence.

Danger zone:

- dashboard becomes decorative
- no failure drilldown
- charts have no denominators

Fallback:

- use Gradio tables and plots only
- skip fancy UI
- focus on metrics and trace inspection

GitHub milestone:

v0.5-dashboard-failure-analysis

## Week 6: Polish, README, Demo, Outreach

Goal:

Make the project understandable to founders in 90 seconds and credible to

engineers in 10 minutes.

Daily deliverables:

|  | Day | Deliverable           | Success condition                |
|--|-----|-----------------------|----------------------------------|
|  | 36  | README rewrite        | clear pitch + architecture       |
|  | 37  | demo GIF              | full loop visible                |
|  | 38  | eval report polish    | metrics table + caveats          |
|  | 39  | code cleanup          | tests/lint pass for core modules |
|  | 40  | interview notes       | answers prepared                 |
|  | 41  | founder outreach list | 30 target startups               |
|  | 42  | final demo rehearsal  | 90-second script works           |

Exact success condition:

Repo is public-ready, demo-ready, and outreach-ready.

Danger zone:

- polishing before demo works
- hiding negative results
- README too long at top

Fallback:

- ship smaller but coherent demo
- write limitations clearly
- emphasize methodology over huge gains

GitHub milestone:

v1.0-internship-demo

## Weeks 7-8: Optional Strengthening

Only if MVP works.

High-value options:

- add Docker sandbox
- expand to 50-100 tasks
- run second DPO experiment
- integrate Phoenix or Langfuse traces
- add AI-assisted failure tagging
- add self-play candidate generation

Do not add:

- browser agents
- Kubernetes
- reward model
- GRPO
- multi-agent workflow

Unless the core loop already works and README is strong.

---

## 68. When You Are Stuck, Do THIS

This section is for the moment when your brain says, "Nothing works." Good. Now debug like an engineer.

### The Stuck Protocol

1. Stop editing.
2. Write the exact command that fails.
3. Run it once.
4. Save the full error.
5. Identify the layer.
6. Reduce to one task.
7. Inspect input and output artifact.
8. Form one hypothesis.
9. Test only that hypothesis.
10. Write the debug note.

### How To Isolate Problems

Ask:

Can I reproduce this with no model?

If yes, it is not an LLM problem.

Examples:

- verifier fails with reference solution -> verifier/task problem
- parser fails on saved raw output -> parser problem
- tool executor fails on manual command -> subprocess/path problem
- preference builder creates wrong pairs -> data pipeline problem
- DPO training crashes before first step -> training/config problem

### Use LLMs Properly For Debugging

Good prompt to an LLM:

I am building a Python CLI. This exact command fails:

<command>

Here is the traceback:

<traceback>

Here is the relevant function:

<function>

I want you to identify the most likely root cause and suggest the smallest patch. Do not rewrite the architecture.

Bad prompt:

My project is broken. Fix it.

Rules:

- give exact error
- give exact code
- ask for minimal patch
- do not paste the whole repo
- verify suggestions manually
- never accept a rewrite unless root cause demands it

## How To Ask Good Technical Questions

Use this format:

Context: I am building X.  
Expected: Y should happen.  
Actual: Z happens.  
Command: ...  
Error: ...  
What I tried: ...  
Smallest relevant code: ...  
Question: Is the issue likely A, B, or C?

This makes you look serious in Discord, GitHub issues, Slack communities, and interviews.

## Avoid Context Overload

When stuck, do not open 20 files. Open:

- the failing command
- the input file
- the component code
- the output/error artifact
- the schema

That is usually enough.

## Avoid Rewriting Everything

Rewrite only when:

- the current abstraction blocks a required feature
- two or more components cannot be tested independently
- data artifacts are impossible to reproduce
- you cannot isolate errors after reasonable effort

Do not rewrite because:

- you feel messy
- a tutorial used a different structure
- a new framework looks cleaner
- you are avoiding a boring bug

## The 30-Minute Rescue Drill

When stuck for more than 30 minutes:

| Minute | Action                                    |
|--------|-------------------------------------------|
| 0-5    | write failing command and expected result |
| 5-10   | run with one task only                    |
| 10-15  | inspect input schema and output artifact  |
| 15-20  | add one debug print/log                   |
| 20-25  | test one hypothesis                       |
| 25-30  | write debug note and decide fix/fallback  |

## Stuck Decision Tree

flowchart TD

```
A["Something failed"] --> B{"Does task validate?"}
B -- "no" --> C["Fix task schema/data"]
B -- "yes" --> D{"Does verifier work manually?"}
D -- "no" --> E["Fix verifier/workspace"]
D -- "yes" --> F{"Does agent produce valid action?"}
F -- "no" --> G["Fix parser/prompt/schema"]
F -- "yes" --> H{"Does tool execute?"}
H -- "no" --> I["Fix subprocess/path/timeout"]
H -- "yes" --> J{"Does trajectory save?"}
J -- "no" --> K["Fix logger/schema"]
J -- "yes" --> L{"Does eval aggregate?"}
L -- "no" --> M["Fix metrics loader"]
L -- "yes" --> N["Inspect model behavior"]
```

## 69. Quality Control Checklists

Use these before README polish, outreach, and interviews.

### Dataset Quality Checklist

- ☐ every task has unique ID
- ☐ broken template fails verifier
- ☐ reference solution passes verifier
- ☐ no train/eval task leakage
- ☐ preference pairs have same prompt for chosen/rejected
- ☐ chosen score > rejected score
- ☐ score margin filter applied
- ☐ random sample manually audited
- ☐ pair count by task reported
- ☐ failure tag distribution reported
- ☐ dataset version frozen

### Verifier Quality Checklist

- ☐ deterministic result
- ☐ timeout enforced
- ☐ stdout/stderr captured
- ☐ exit code captured
- ☐ failure tags assigned
- ☐ reference pass checked
- ☐ broken fail checked
- ☐ protected files considered
- ☐ false positive examples tested
- ☐ false negative examples investigated

### Evaluation Credibility Checklist

- ☐ held-out task split used
- ☐ same config for base/tuned
- ☐ same verifier for base/tuned
- ☐ same max steps/timeouts
- ☐ denominator reported



- ☐ pass rate reported
- ☐ invalid action rate reported
- ☐ timeout rate reported
- ☐ failure tags reported
- ☐ limitations written

## Dashboard Quality Checklist

- ☐ loads real run artifacts
- ☐ shows base vs tuned metrics
- ☐ shows failure distribution
- ☐ shows trace viewer
- ☐ shows chosen/rejected pair examples
- ☐ no fake placeholder metrics
- ☐ charts have labels and denominators
- ☐ errors are visible if files missing

## README Quality Checklist

- ☐ one-line pitch clear
- ☐ demo GIF or screenshot above fold
- ☐ architecture diagram present
- ☐ results table present
- ☐ setup commands present
- ☐ eval methodology explained
- ☐ failure analysis included
- ☐ limitations honest
- ☐ tech stack listed
- ☐ roadmap shows judgment

## Demo Quality Checklist

- ☐ 90-second script rehearsed
- ☐ one successful run ready
- ☐ one failure analysis example ready
- ☐ dashboard loads quickly
- ☐ commands are reproducible

- ☐ no secret/API key shown
- ☐ metrics match README
- ☐ fallback video/GIF exists

## Recruiter-Readability Checklist

- ☐ title says what project does
- ☐ first paragraph explains business relevance
- ☐ screenshot shows real system
- ☐ tech stack matches AI infra/post-training/evals roles
- ☐ resume bullets include metrics
- ☐ GitHub repo is not cluttered
- ☐ README does not require reading a paper first

## Final Pre-Outreach Checklist

Can a founder understand the project in 30 seconds?  
 Can an engineer inspect the repo in 10 minutes?  
 Can you demo the full loop in 90 seconds?  
 Can you answer what failed?  
 Can you prove the metrics came from held-out evals?

If yes, start outreach.

---

## 70. Final Daily Navigator

When you sit down tomorrow, do not reread the whole handbook.

Do this:

1. Open Section 55.
2. Write today's one artifact.
3. Run the first command.
4. If stuck, open Section 68.
5. If implementing a component, open Section 56.
6. End with a note, artifact, and commit.

The strongest version of this project is not the largest version. It is the version where every claim has an artifact behind it:

claim -> command -> artifact -> metric -> explanation

That is what makes it internship-grade.